



# Kubernetes on Azure 101 Workshop

Rui Félix Pereira

2020.05.05




# Containers & Kubernetes workshops

## Containers 101



## Kubernetes 101



## Kubernetes 201



# Introduction to Kubernetes

# Containers are great but “at scale” become very challenging



Standardize Build




DevOps CI / CD



Run Anywhere




Rapid deployment



Compute Density & Scale

Deploy, manage and update 1000s of containers in production



# Container management & orchestration “at scale”



Scheduling



Affinity  
anti-affinity




Health  
monitoring



Failover



Scaling



Networking



Service  
discovery



Coordinated  
app upgrades

# Kubernetes: the industry leading orchestrator



## Portable

Public, private, hybrid,  
multi-cloud



## Extensible

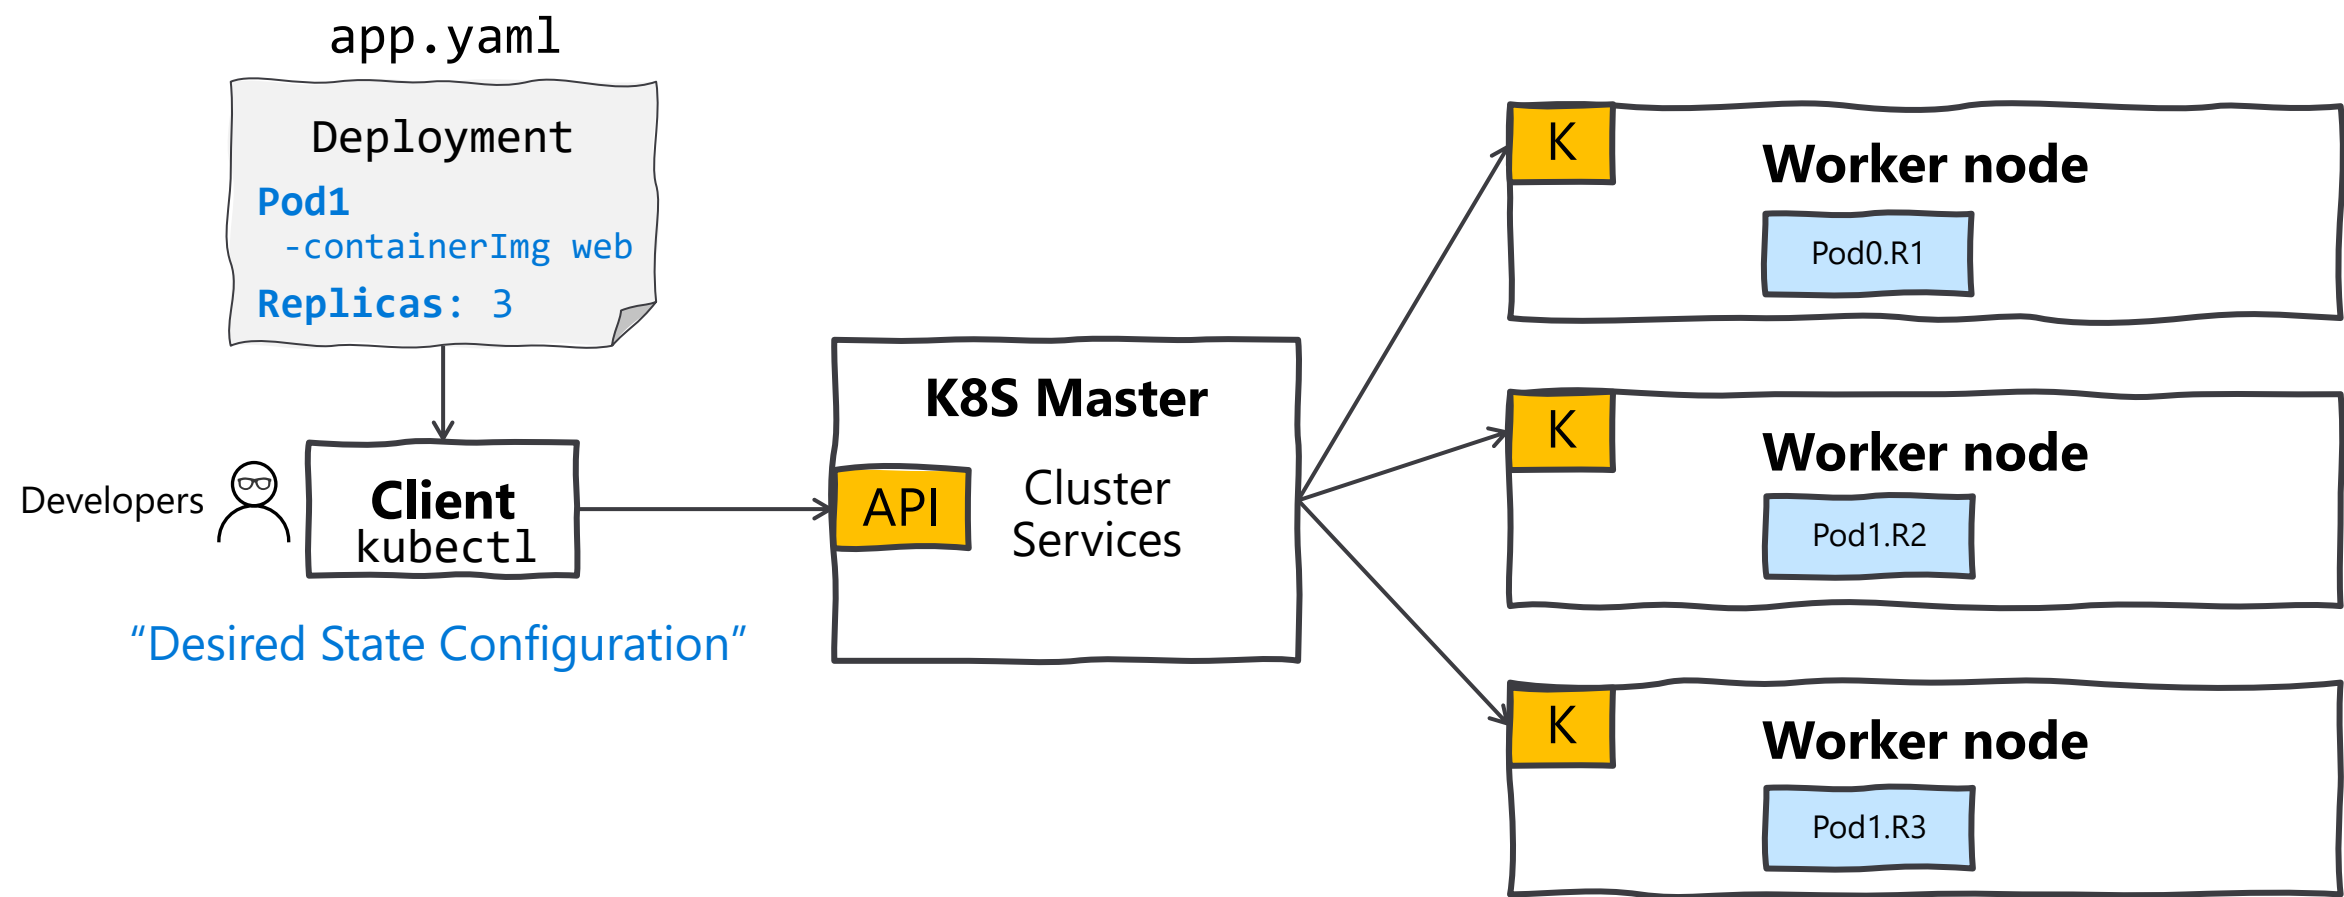
Modular, pluggable,  
hookable, composable



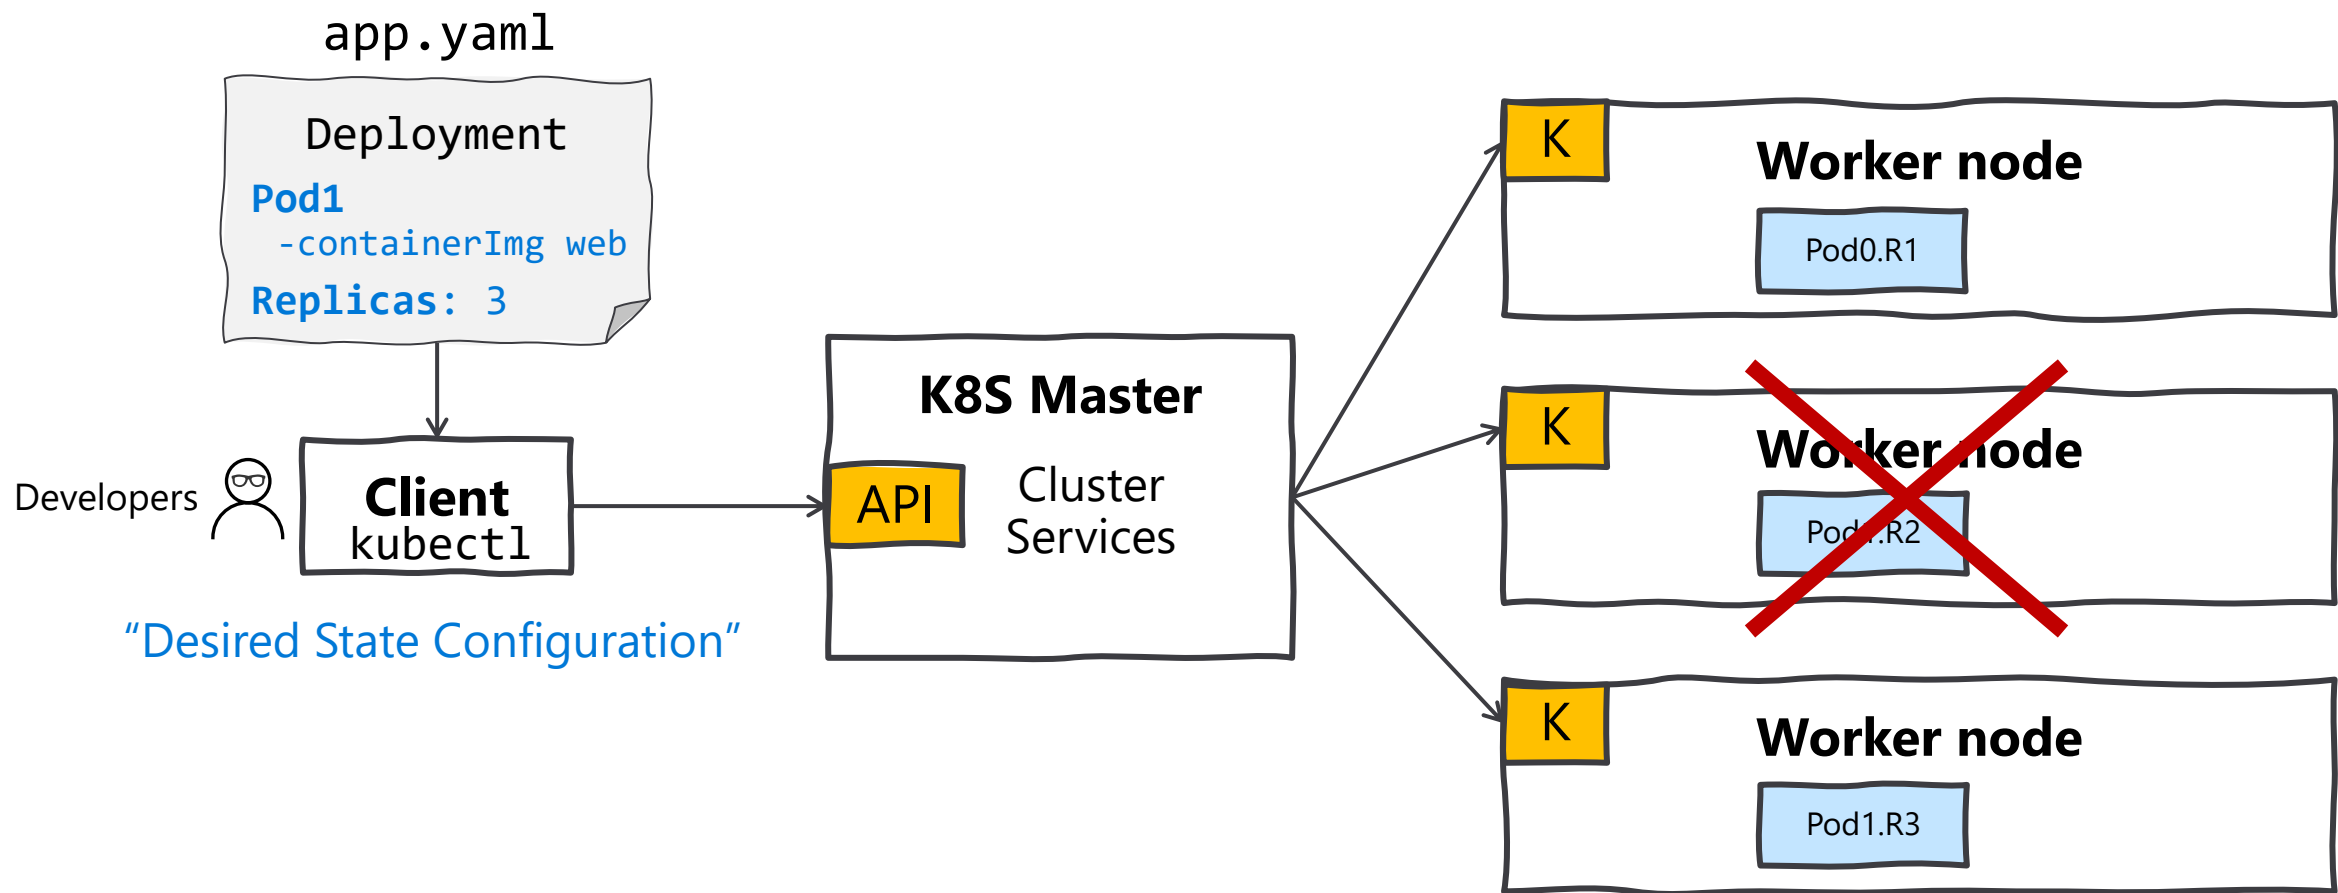
## Self-healing

Auto-placement, auto-restart,  
auto-replication, auto-scaling

# Kubernetes 101

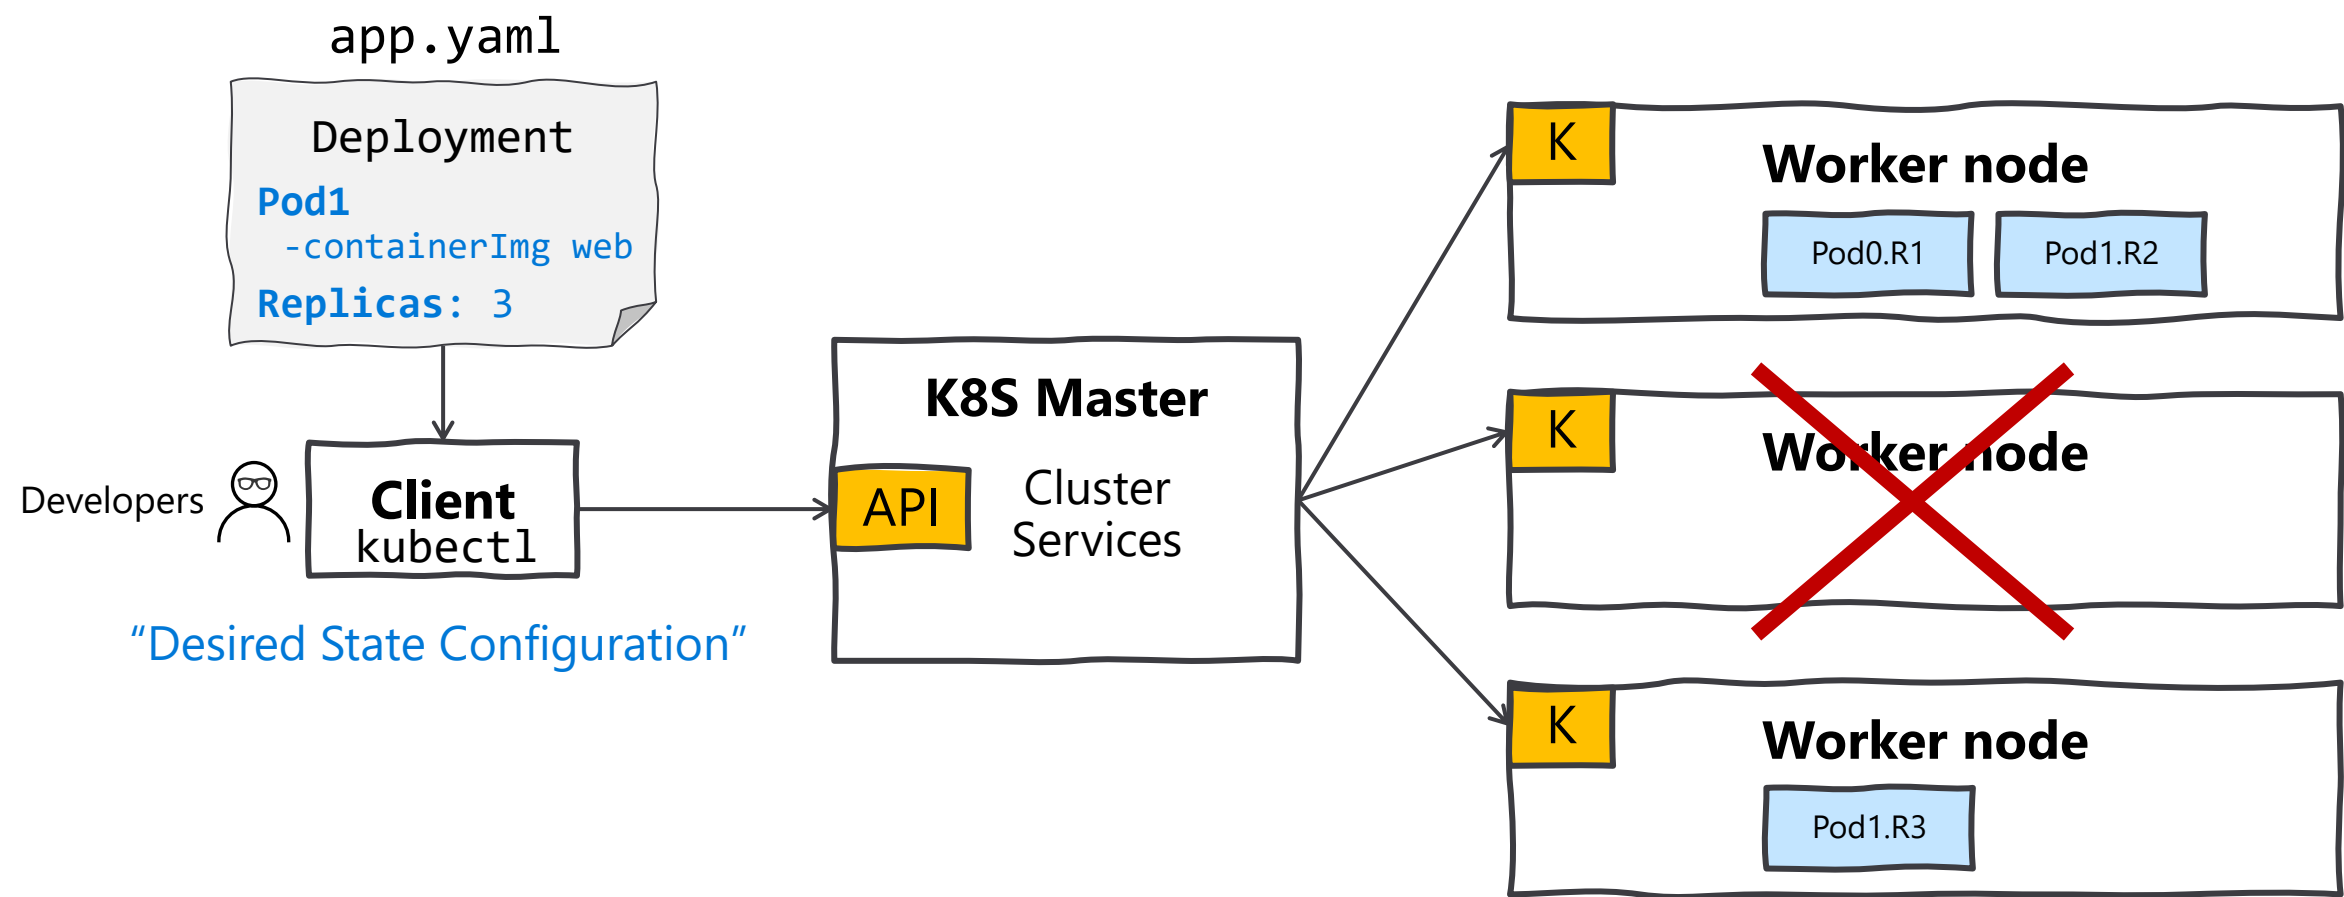


# Kubernetes 101

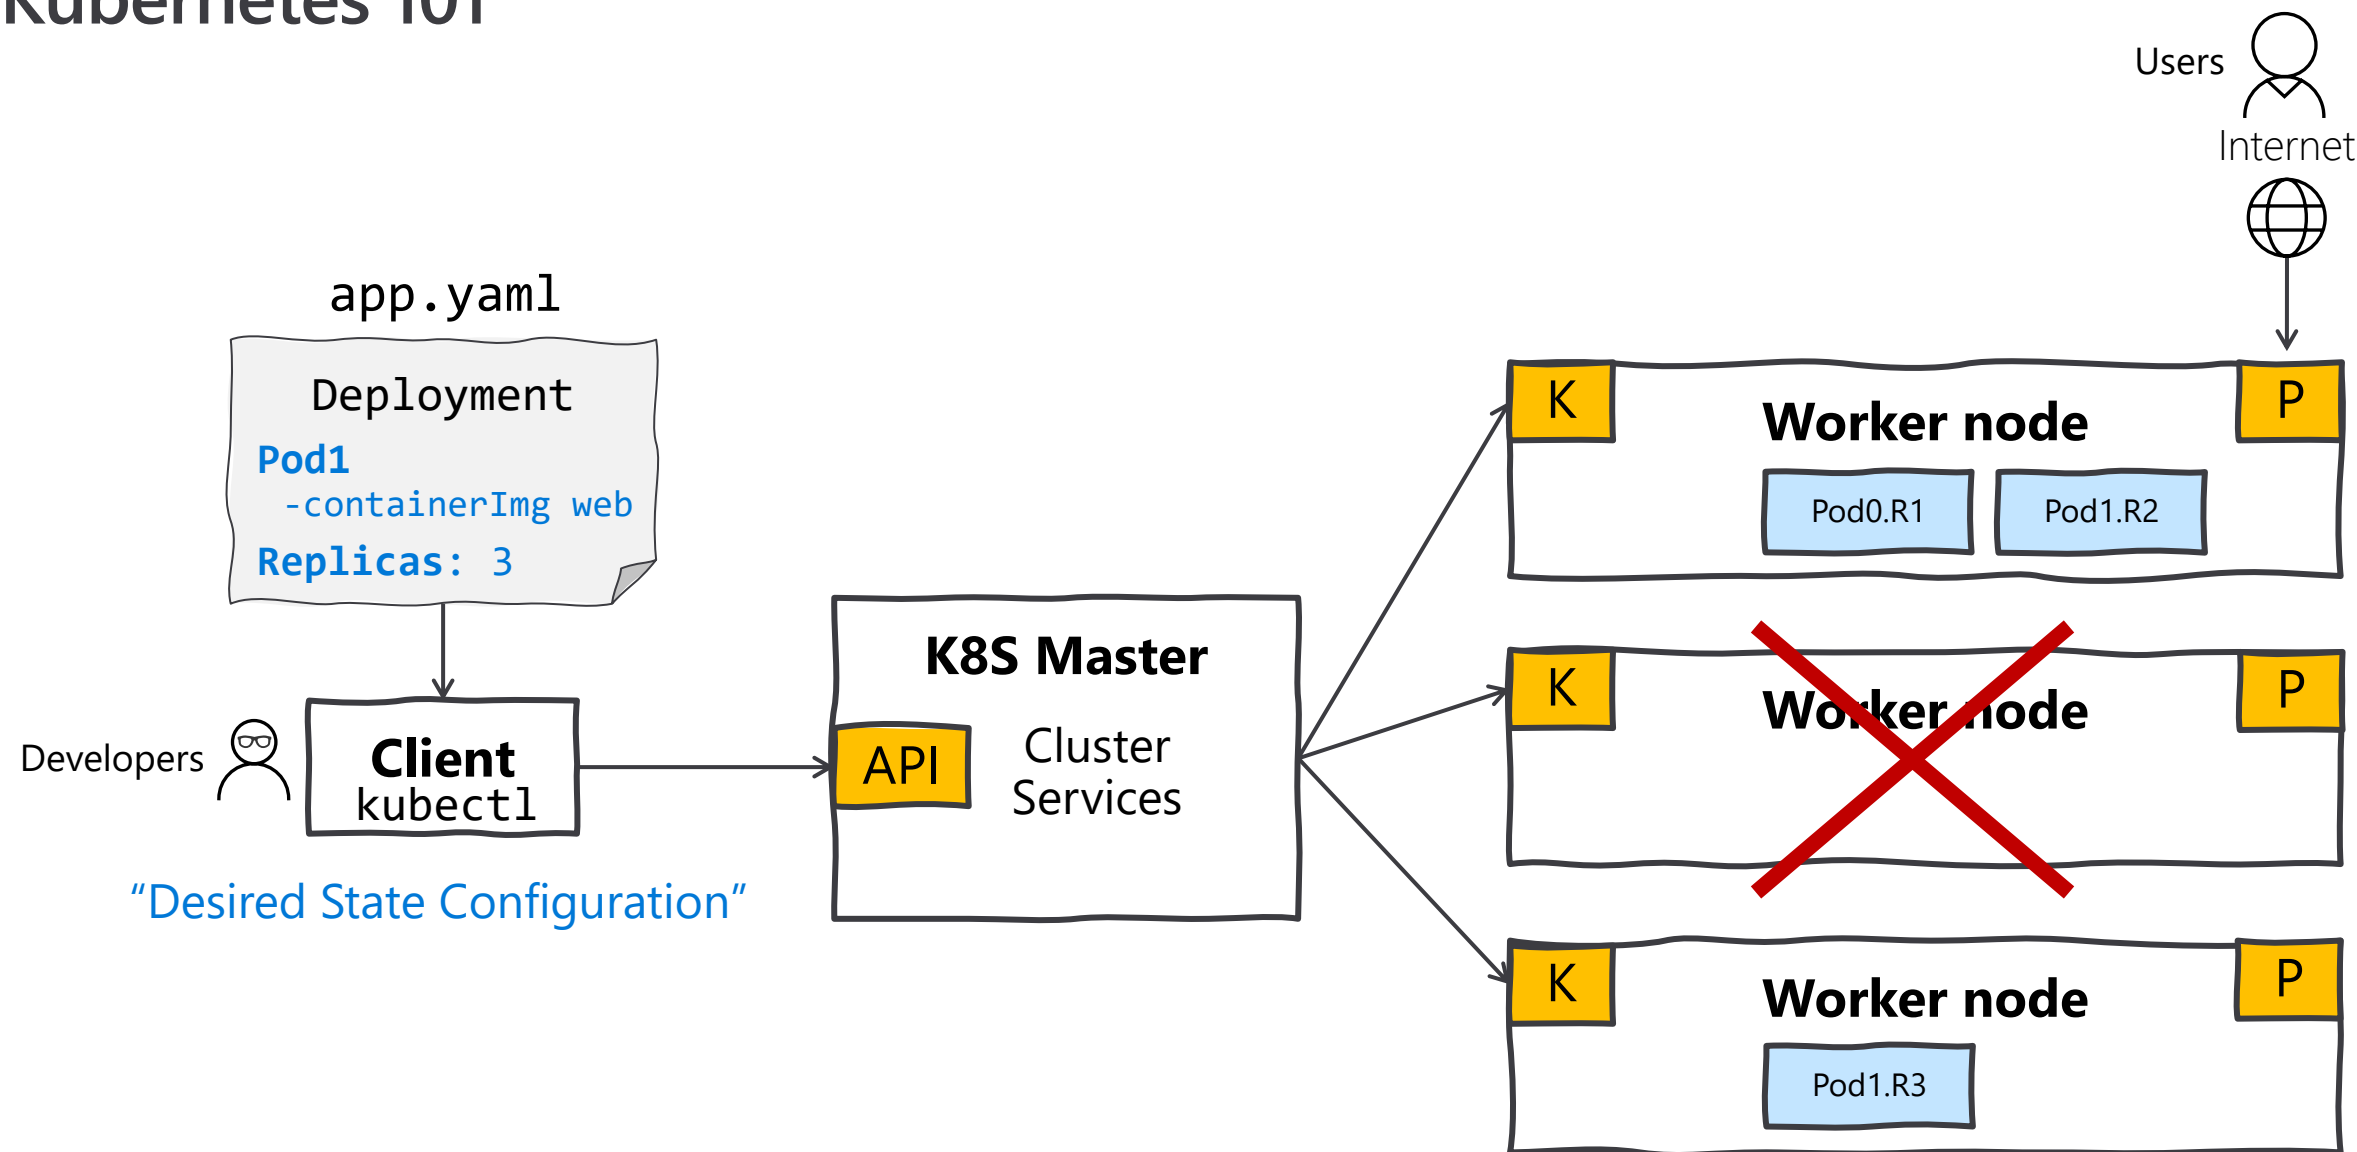




# Kubernetes 101




# Kubernetes 101




# Kubernetes on Azure

# Azure Kubernetes Service (AKS)


Simplify deployment, management and operations of Kubernetes




Manage Kubernetes  
with ease



Accelerate  
containerized  
development



Build on an  
enterprise-grade,  
secure foundation



Run anything,  
anywhere

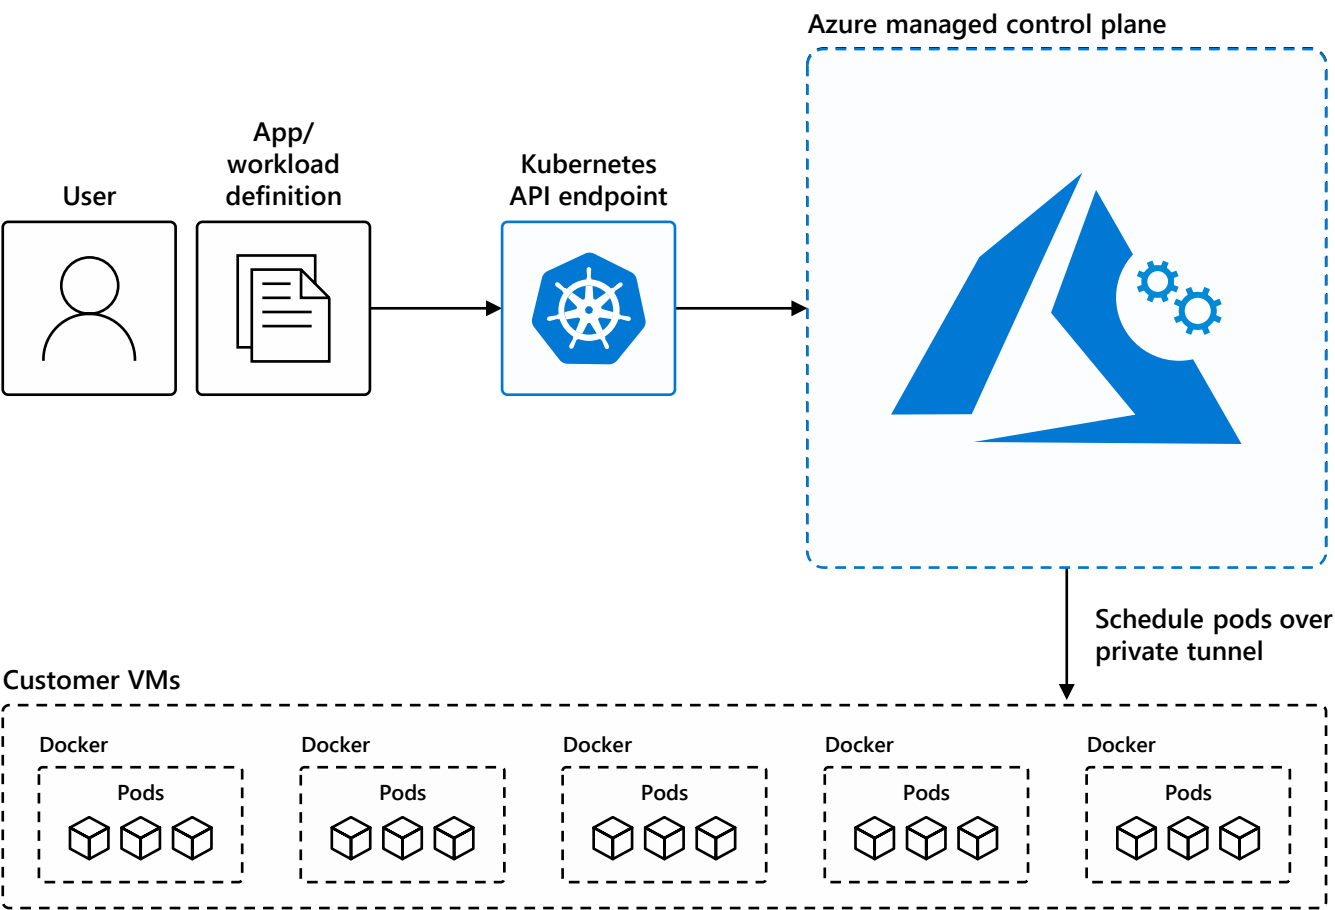


# Increase operational efficiency

Focus on your containers and code, not the plumbing of them

Responsibilities	DIY with Kubernetes	Managed Kubernetes on Azure
Containerization	<div></div>	<div></div>
Application iteration, debugging	<div></div>	<div></div>
CI/CD	<div></div>	<div></div>
Provisioning, upgrades, patches	<div></div>	<div></div>
Reliability availability	<div></div>	<div></div>
Scaling	<div></div>	<div></div>
Monitoring and logging	<div></div>	<div></div>

Customer Microsoft

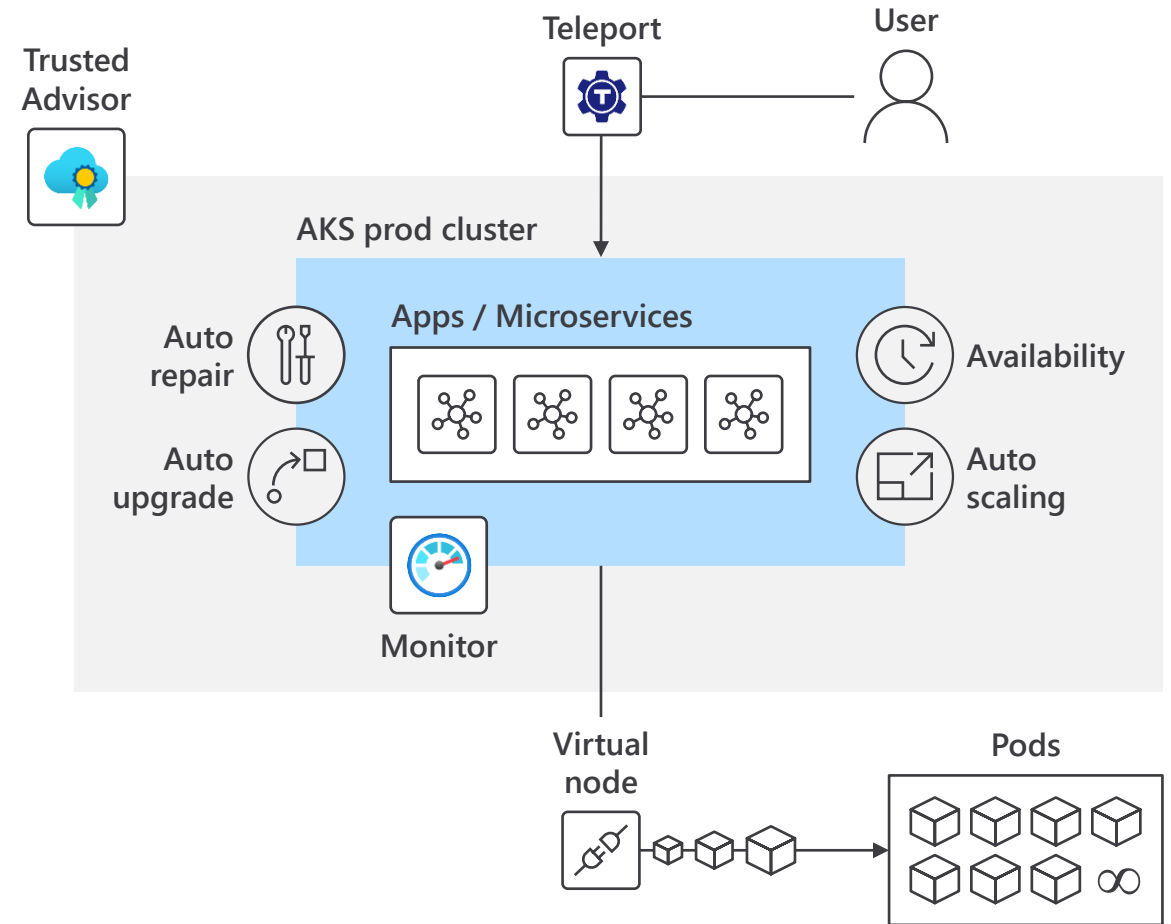




# Increased operational efficiency


Highly available, reliable service with serverless scaling

- Easily provision fully managed clusters with automatically configured monitoring capabilities based on Prometheus
- Real-time personalized recommendations to optimize your AKS deployments with Azure Advisor integration
- Elastically add compute capacity with serverless Kubernetes in seconds without worrying about managing the infrastructure.
- Higher availability using redundancies across availability zones, protecting applications from datacenter failures



# Serverless Kubernetes using AKS virtual nodes

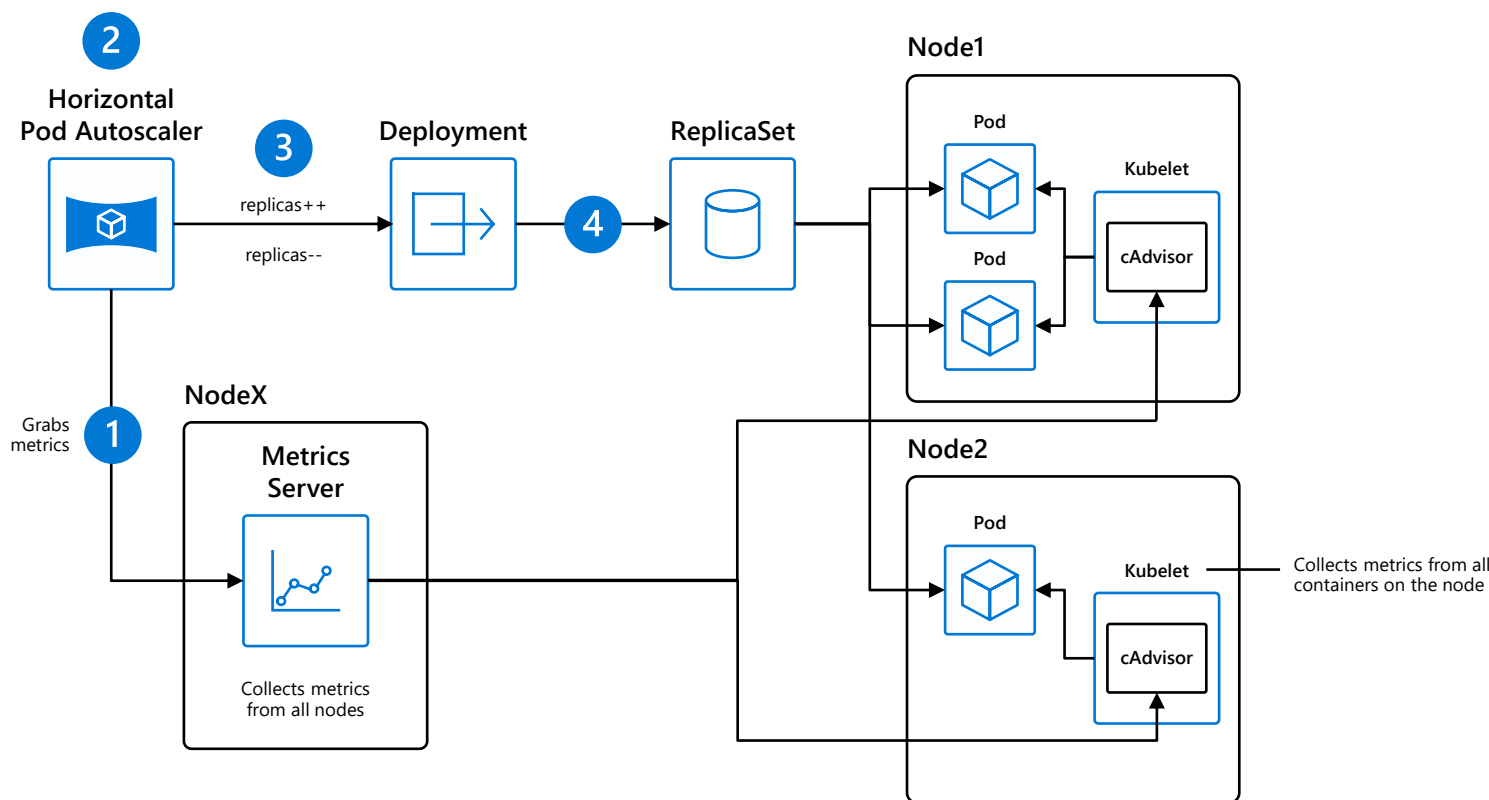
- Elastically provision compute capacity in seconds
- No infrastructure to manage
- Built on open sourced Virtual Kubelet technology, donated to the Cloud Native Computing Foundation (CNCF)



# Horizontal Pod Autoscaler

The horizontal pod autoscaler (HPA) uses the Metrics Server in a Kubernetes cluster to monitor the resource demand of pods. If a service needs more resources, the number of pods is automatically increased to meet the demand.

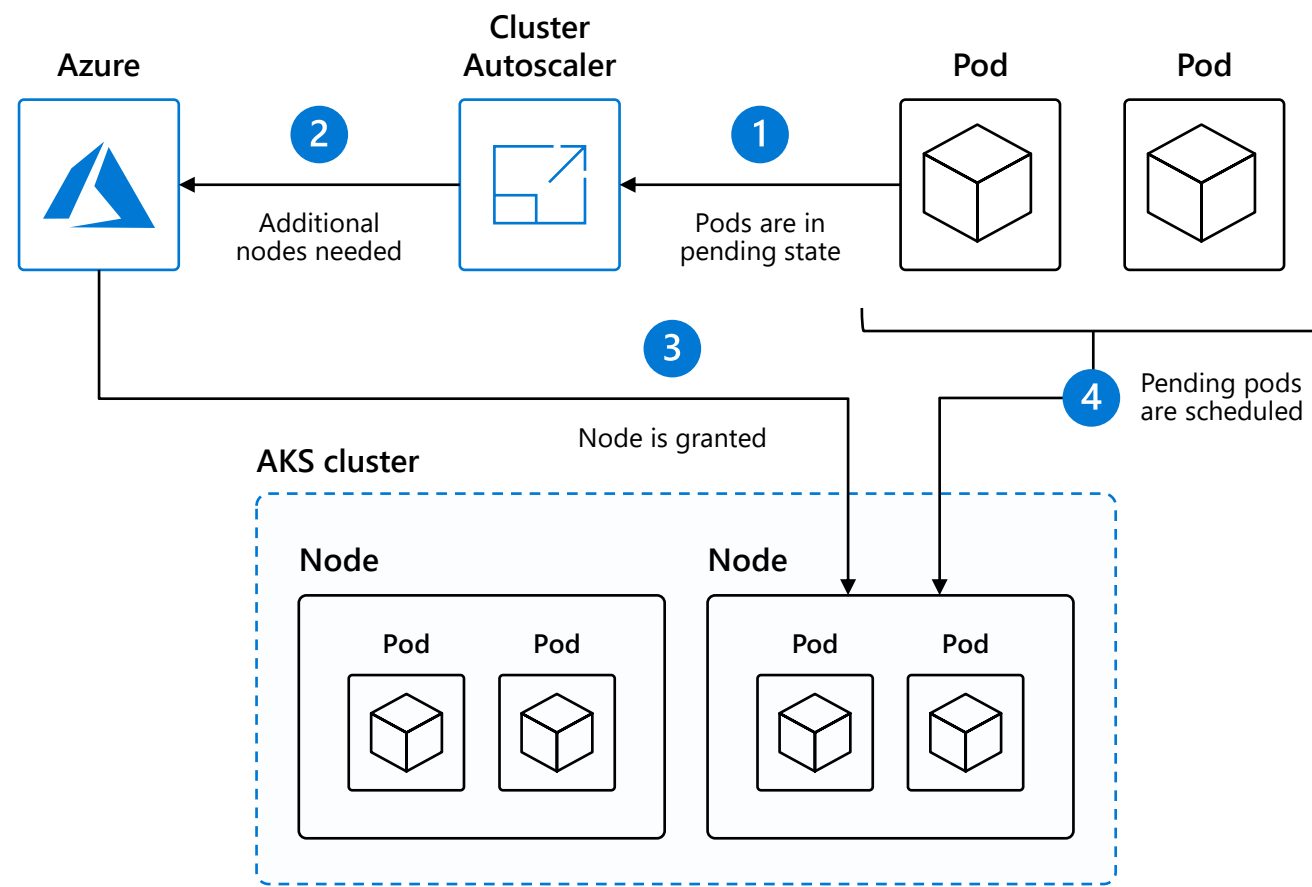
1. HPA obtains resource metrics and compares them to user-specified threshold
2. HPA evaluates whether user specified threshold is met or not
3. HPA increases/decreases the replicas based on the specified threshold
4. The Deployment controller adjusts the deployment based on increase/decrease in replicas




# Cluster Autoscaler

*The cluster autoscaler watches for pods that can't be scheduled on nodes because of resource constraints. The cluster then automatically increases the number of nodes.*

1. HPA obtains resource metrics and compares them to user-specified threshold
2. HPA evaluates whether user specified threshold is met or not
3. HPA increases/decreases the replicas based on the specified threshold
4. The Deployment controller adjusts the deployment based on increase/decrease in replicas



# Build on an enterprise-grade, secure platform



Control  
access through  
AAD and RBAC



Get runtime vulnerability  
scanning and auditing  
through Azure Security Center



Put guardrails in your  
development process with  
Azure Policy



Secure network  
communications with VNET  
and network policy



Gain automated threat  
protection and best practice  
recommendations for  
Kubernetes clusters

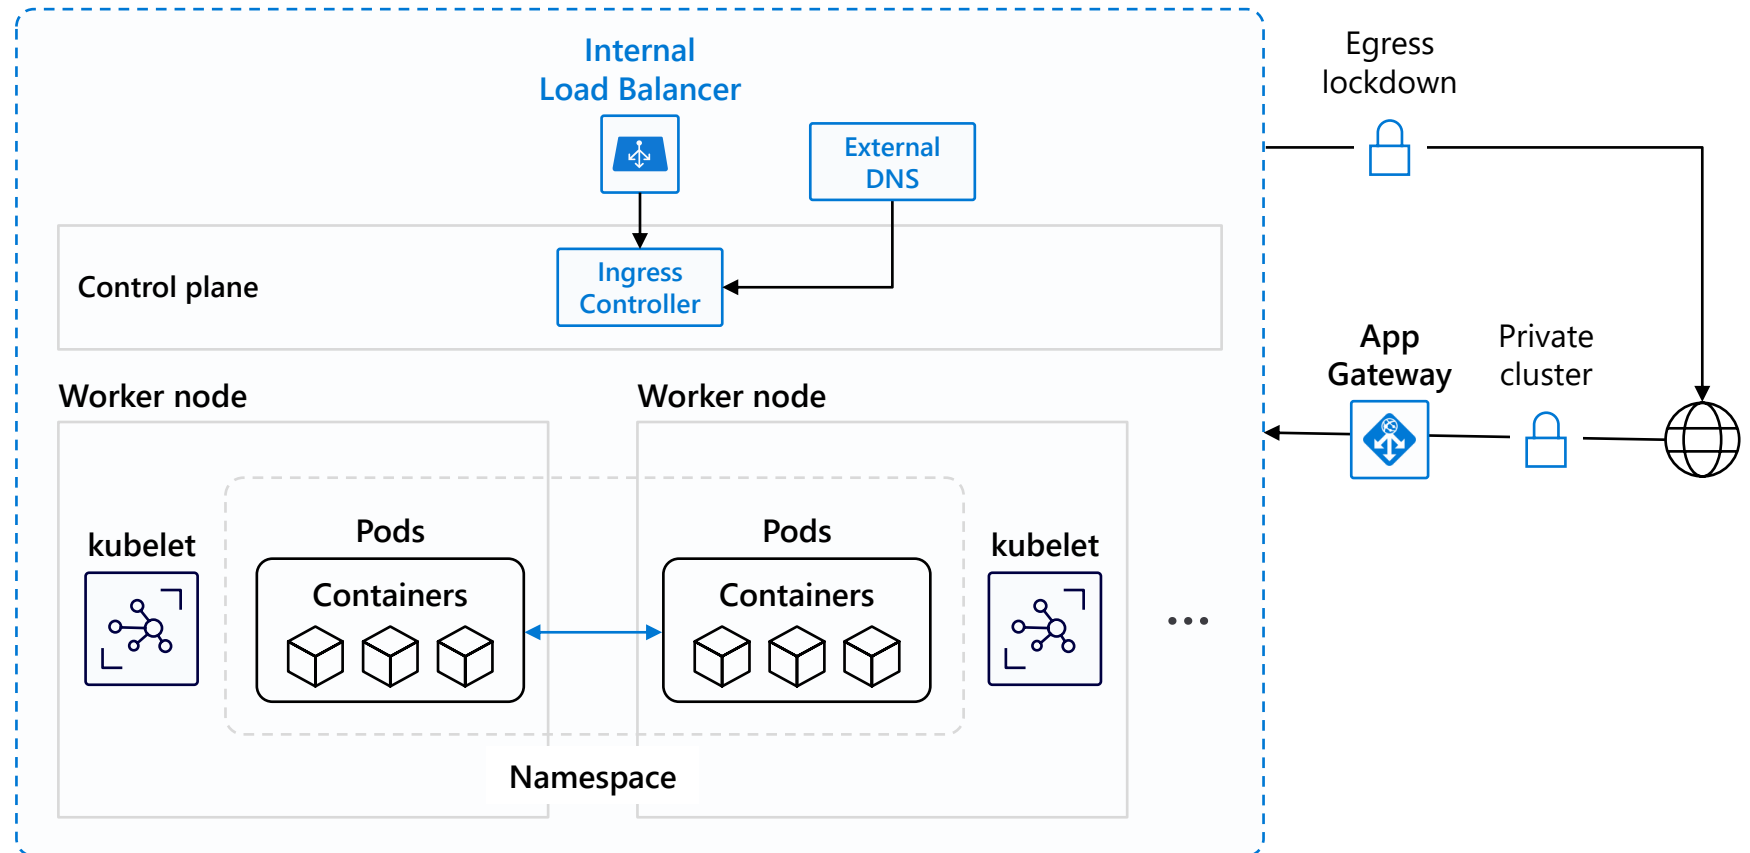




# Networking

Secure your Kubernetes workloads with [virtual network](#) and policy-driven communication paths between resources

Kubernetes cluster: [Azure VNET](#)



# Basic networking

## Uses kubenet network plugin and has the following features

Nodes and Pods are placed on different IP subnets

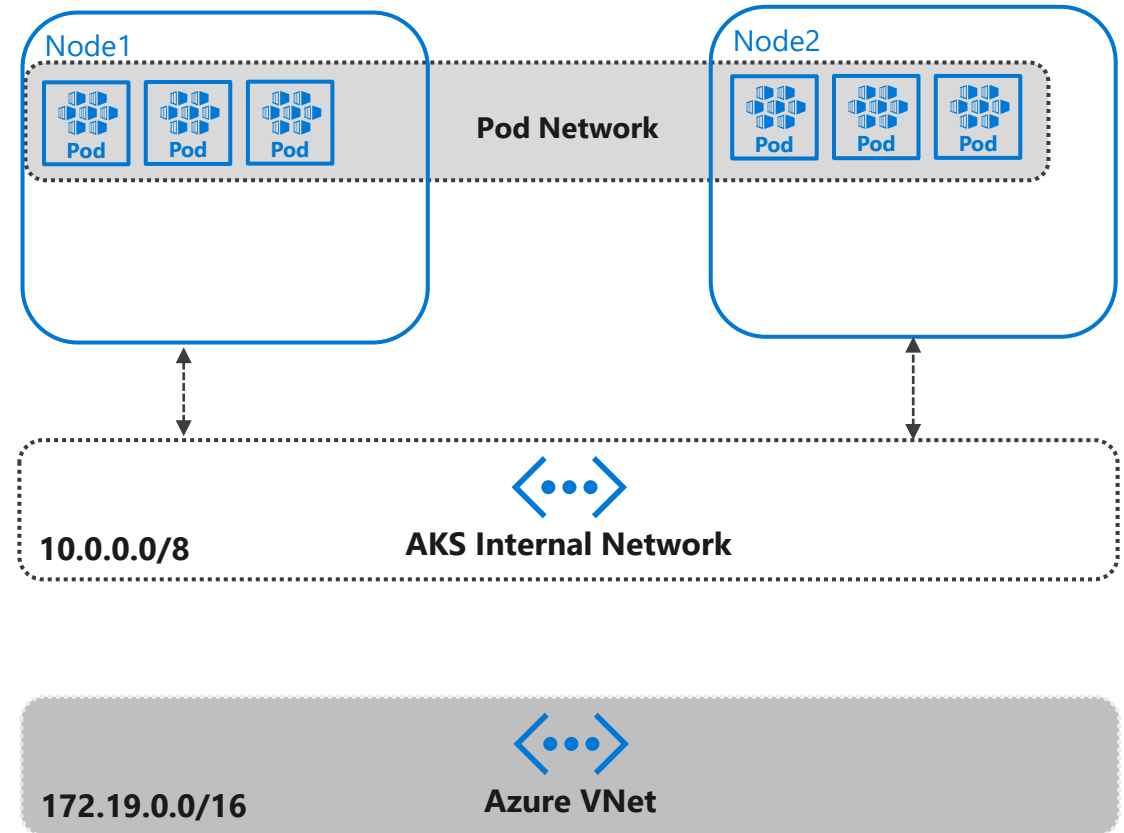
User Defined Routing and IP Forwarding is for connectivity between Pods across Nodes.

## Drawbacks

2 different IP CIDRs to manage

Performance impact

Peering or On-Premise connectivity is hard to achieve



# Advanced networking

## Uses the Azure CNI (Container Networking Interface)

CNI is a vendor-neutral protocol, used by container runtimes to make requests to Networking Providers

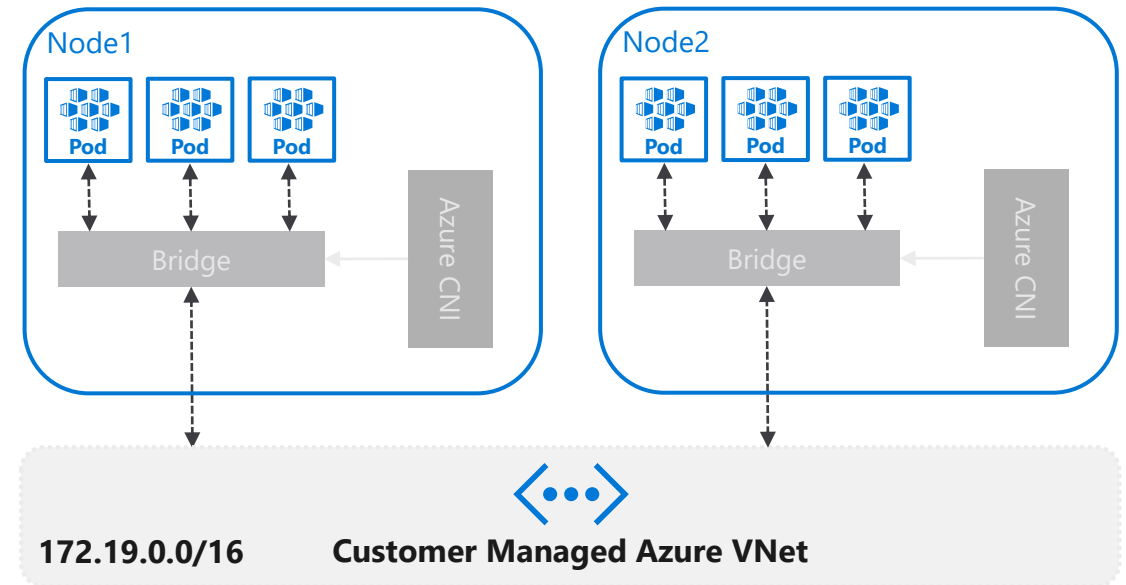
Azure CNI is an implementation which allows you to integrate Kubernetes with your VNET

## Advantages


Single IP CIDR to manage

Better Performance

Peering and On-Premise connectivity is out of the box



# Advanced networking: Planning IP addressing for your cluster



## Additional subnets

### Kubernetes service address range

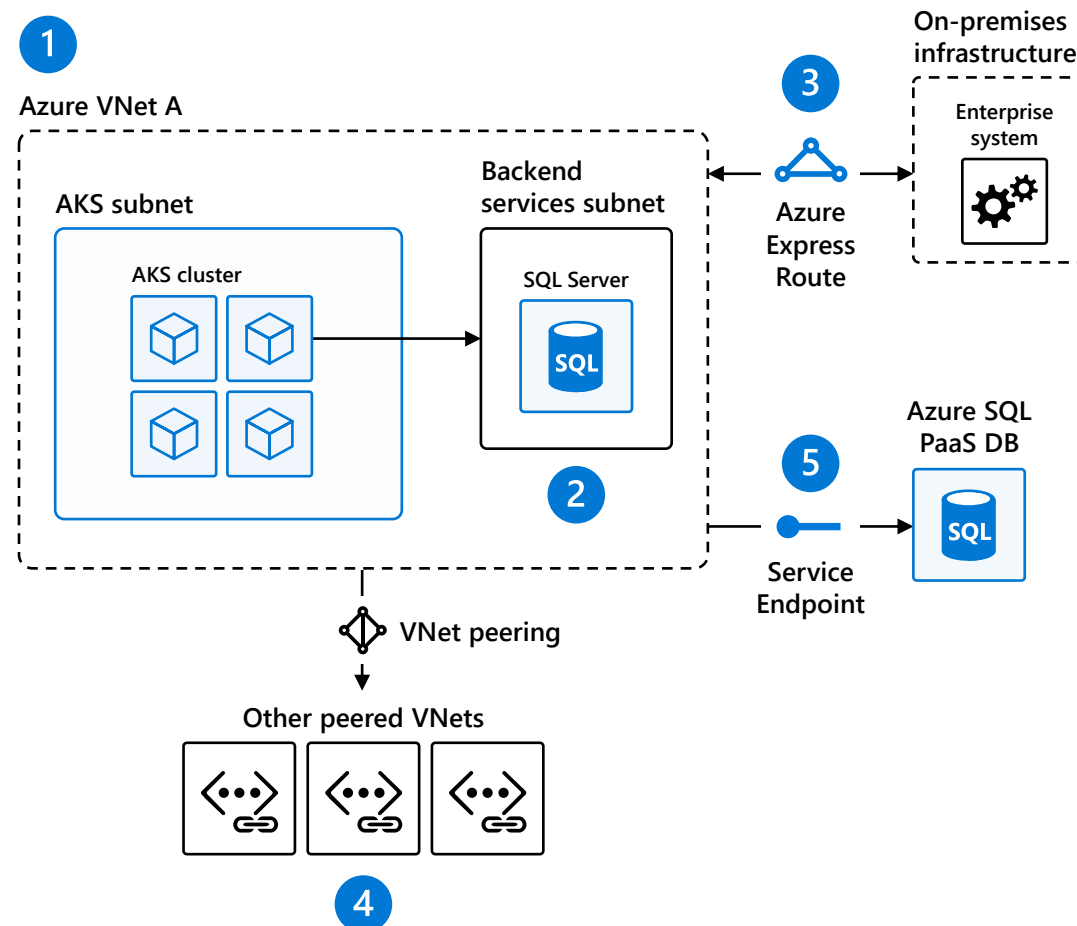
- Non-overlapping with other subnets in your network (does not need to be routable) and not 169.254.0.0/16, 172.30.0.0/16 and 172.31.0.0/16.
- Smaller than /12

### Docker bridge address

- Non-overlapping with AKS Subnet

# Secure network communications with VNET and CNI

1. Uses Azure subnet for both your containers and cluster VMs
2. Allows for connectivity to existing Azure services in the same VNet
3. Use Express Route to connect to on-premises infrastructure
4. Use VNet peering to connect to other VNets
5. Connect AKS cluster securely and privately to other Azure resources using VNet endpoints




AKS VNet integration works seamlessly with your existing network infrastructure



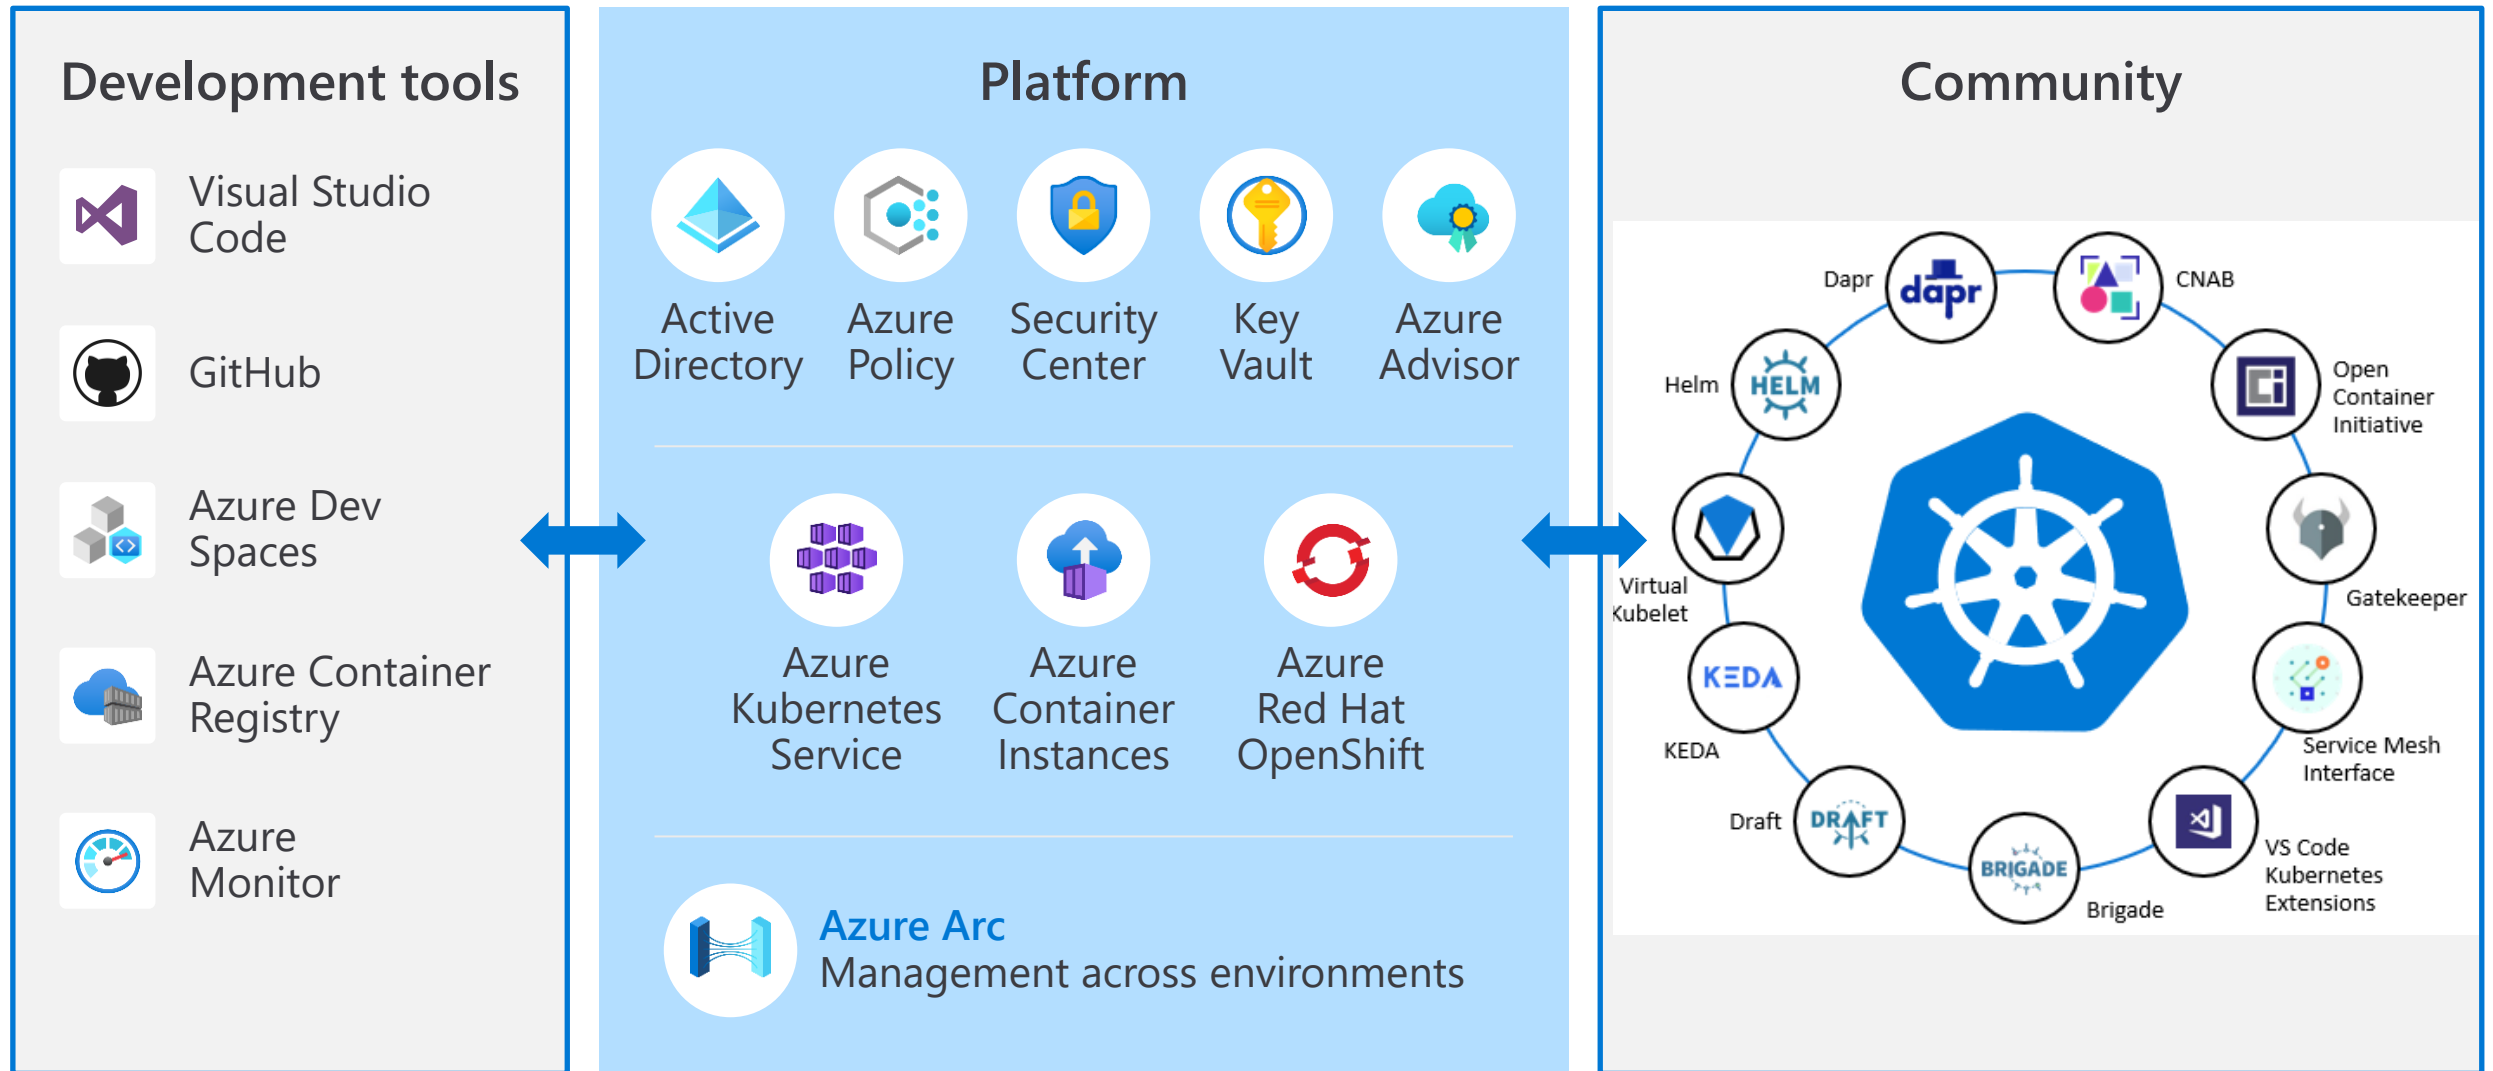
# Azure Kubernetes Service (AKS) support for Windows Server Containers

Now you can get the best of managed Kubernetes for all your workloads whether they're in Windows, Linux, or both

- Lift and shift Windows applications to run on AKS
- Seamlessly manage Windows and Linux applications through a single unified API
- Mix Windows and Linux applications in the same Kubernetes cluster—with consistent monitoring experience and deployment pipelines




# Kubernetes on Azure | Enterprise-grade by design



# Unified management

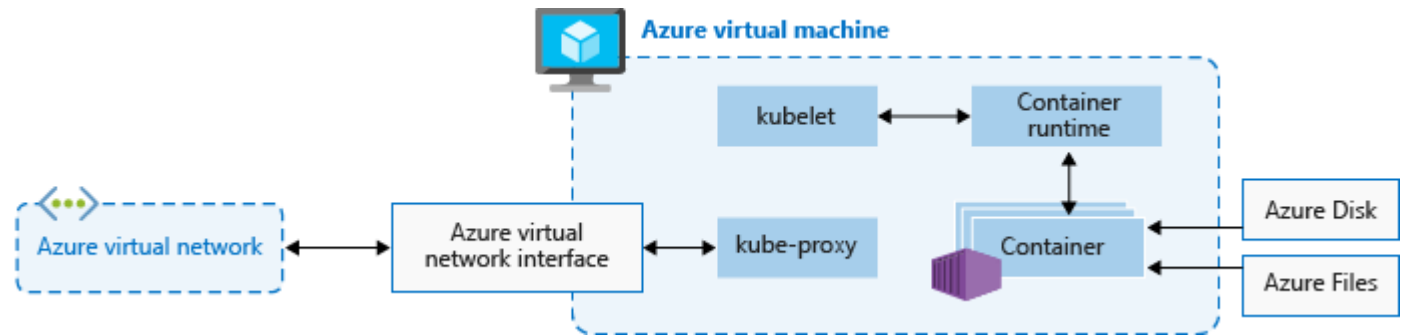
- Central inventory and monitoring of the sprawling assets running anywhere from on-premises to edge
- Consistently apply policies, role-based-access-controls (RBAC) for at-scale governance
- Deploy Kubernetes resources to all clusters using a GitOps-based workflow



# Create your first Kubernetes cluster

# AKS: Nodes and node pools

- AKS can have multiple node pools (Linux and Windows)
- Each node is an Azure VM
- Azure VM size for your nodes defines how many CPUs, how much memory, and the size and type of storage available (such as high-performance SSD or regular HDD)





# Lab 0: Setup

Check proper access and customize your environment



Task



With Azure

Have access to Azure Subscription for AKS deployment

Check if you have access to an Azure Subscription for deploying your own AKS managed cluster.  
Test access to Azure Cloud Shell.

Create or use a Service Principal

Create a new Service Principal or get one, namely the service principal application client ID and the corresponding client secret. This will be used for AKS deployment.

Copy .env file

```
mv .env.customize .env
```

Customize .env

Edit your .env file and customize the following properties:

```
LOCATION=northeurope
APP_NAME=azure-vote
RESOURCE_GROUP=k8s101
CLUSTER_NAME=k8s101
SUBSCRIPTION=<your subscription id>
SERVICE_PRINCIPAL_APP_ID=<your service principal application client id>
SERVICE_PRINCIPAL_SECRET=<your service principal client secret>
```

# Lab 1: Nodes and node pools

Create a managed Azure Kubernetes Service cluster



Task



With Azure

Create a cluster

```
az group create --name $(RESOURCE_GROUP) --location $(LOCATION) --output table
az aks create \
    --resource-group $(RESOURCE_GROUP) \
    --name $(CLUSTER_NAME) \
    --node-count 2 \
    --node-vm-size Standard_D4s_v3 \
    --generate-ssh-keys
--no-wait
```

Validate cluster creation

```
sudo az aks install-cli          # not needed if running on Azure cloud shell
az aks show \
    --resource-group $(RESOURCE_GROUP) \
    --name $(CLUSTER_NAME)
az aks get-credentials \
    --resource-group $(RESOURCE_GROUP) \
    --name $(CLUSTER_NAME)

kubectl get nodes
kubectl cluster-info
```

# Lab 1b: Kubectl alias

Just a suggestion



Task



With Azure

Create kubectl alias

# Edit your `.bashrc` and add this line:

```
alias k='kubectl'
```

# Apply the new settings

```
source .bashrc
```

# Test it

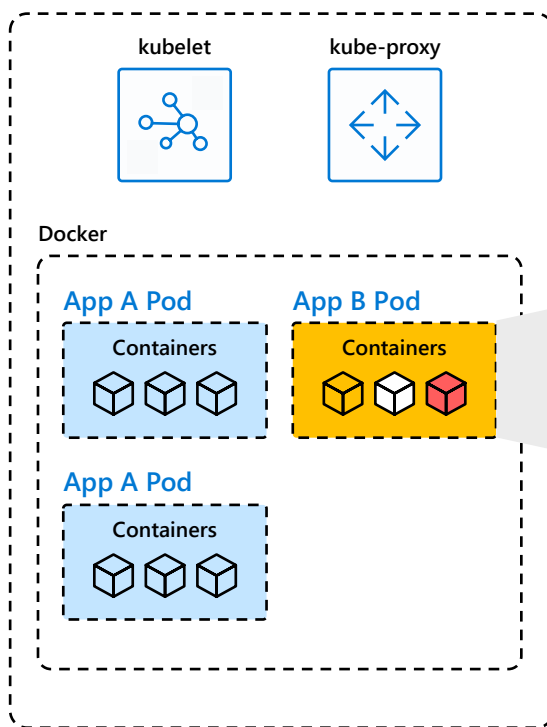
```
k get nodes
```

# Pods & Deployments core concepts


# Pods: The building block for your applications in K8S

- Encapsulates an application container (or multiple containers), storage resources, a unique network IP, and options that govern how the container(s) should run
- Pods group containers that work together
- Represents a unit of deployment and scaling: a single instance of an application in Kubernetes

## Worker node




## App B Pod



# Pods - Single container deployment

- Pods that run a single container.
- The “one-container-per-Pod” model is the most common Kubernetes use case
- You can think of a Pod as a wrapper around a single container, and Kubernetes manages the Pods rather than the containers directly


App C Pod



```
apiVersion: v1
kind: Pod
metadata:
  name: helloWeb
spec:
  containers:
    image: nginx
    ports:
      - containerPort: 80
```

# Pods - Multiple containers deployment

## Web App Pod

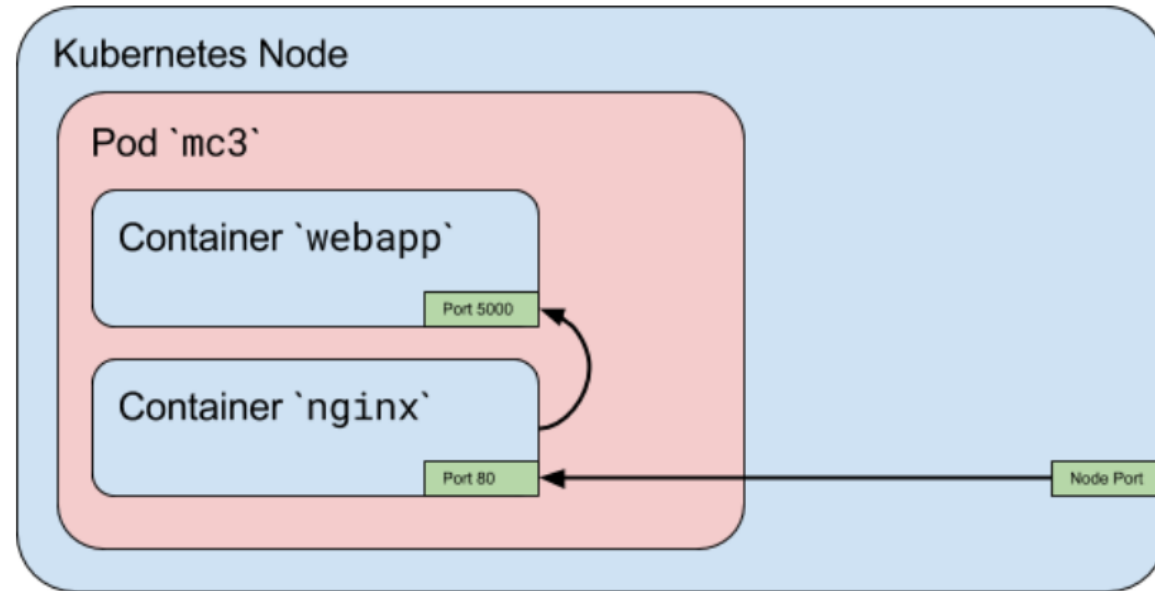


- Pods that run multiple containers that need to work together
- A Pod might encapsulate an application composed of **multiple co-located containers** that are tightly coupled and **need to share resources**

```
apiVersion: v1
kind: Pod
spec:
  volumes:
    - name: html
      emptyDir: {}
  containers:
    - name: webserver
      image: nginx
      volumeMounts:
        - name: html
          mountPath: /usr/share/nginx/html
    - name: content
      image: debian
      volumeMounts:
        - name: html
          mountPath: /html
```

# Pods - Network

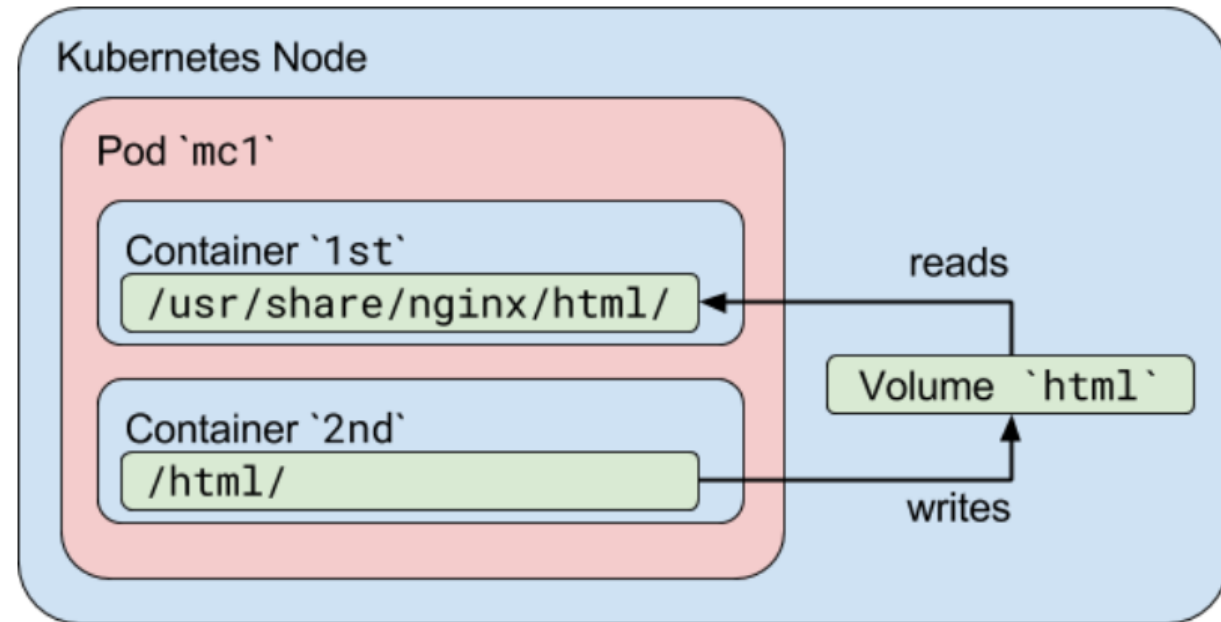
- Each Pod is assigned a unique IP address
- Every container in a Pod shares the network namespace, including the IP address and network ports
- Containers inside a Pod can communicate with one another using localhost
- When containers in a Pod communicate with entities outside the Pod, they must coordinate how they use the shared network resources (such as ports)





# Pods - Storage

- A Pod can specify a set of shared storage volumes
- All containers in the Pod can access the shared volumes, allowing those containers to share data
- Volumes also allow persistent data in a Pod to survive in case one of the containers within needs to be restarted



# Pods – Resource limits

- When Containers have resource requests specified, the scheduler can make better decisions about which nodes to place Pods on
- When Containers have their limits specified, contention for resources on a node can be handled in a specified manner

```
apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
  - name: db
    image: mysql
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
```

# Deployments for fault tolerance and scale

- A **Deployment** controller provides declarative updates for **Pods** and **ReplicaSets**
- You describe a desired state in a Deployment object, and the Deployment controller changes the actual state to the desired state at a controlled rate
- **How it works:**
  - Deployment controller creates a ReplicaSet
  - Deployment controller scales the ReplicaSet to the desired count
  - ReplicaSet creates the desired number of Pods

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-world-deployment
  labels:
    app: hello-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hello-app
  spec:
    containers:
      - name: hello-world
        image: hello
```

**Deploy & scale  
your app**

# Lab 2: Deployments

Deploy sample application and scale it



Task



With Azure

Deploy sample app

```
kubectl create -f azure-vote.yaml
```

Validate new app

```
# get-deployments  
kubectl get deployments
```

```
# describe-deployments  
kubectl describe deployments
```

```
# get-replicaset  
kubectl get rs
```

```
# get-pods  
kubectl get pods --show-labels
```

Scale your app

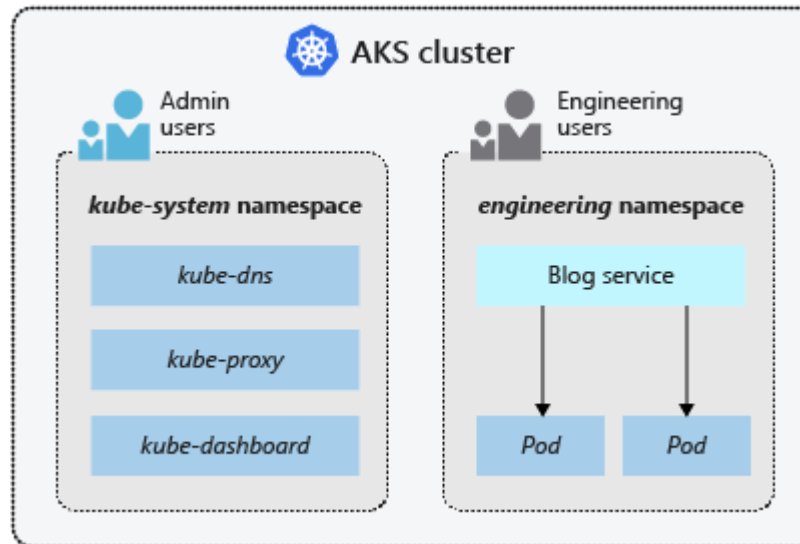
```
kubectl scale deployment azure-vote-front --replicas=2
```

# Organize with Namespaces

# Namespaces

- Namespaces provide a scope for names. Names of resources need to be unique within a namespace
- Namespaces are a way to divide cluster resources between multiple users (via resource quota)
- A service DNS entry is of the form `<service-name>.<namespace-name>.svc.cluster.local`, which means that if a container just uses `<service-name>`, it will resolve to the service which is local to a namespace

```
$ kubectl --namespace=<namespace-name> get pods  
  
$ kubectl config set-context $(kubectl config current-context) --namespace=<namespace-name>
```



```
apiVersion: v1  
kind: Namespace  
metadata:  
  name: engineering
```

# Lab 3: Namespaces

## Create new namespace



Task



With Azure

Check existing namespaces

```
kubectl get namespaces
```

Create namespace

```
kubectl create namespace hellons
```

Check new namespace

```
kubectl get namespaces
```

```
# get pods in the new namespace
```

```
kubectl --namespace=hellons get pods
```



# Expose your app using Services


# Services to expose and load balance applications

- Services logically group a set of pods together and provide network connectivity
- Allow to expose endpoints (public or private)
- Load balance requests between app replicas
- Several Service types are available

```
kind: Service
metadata:
  name: azure-vote-front
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: azure-vote-front
```

# Services – Cluster IP


- Creates an internal IP address for use within the AKS cluster
- Good for internal-only applications that support other workloads within the cluster



```
kind: Service
metadata:
  name: azure-vote-back
spec:
  ports:
    - port: 80
  selector:
    app: azure-vote-back
```

# Services – NodePort

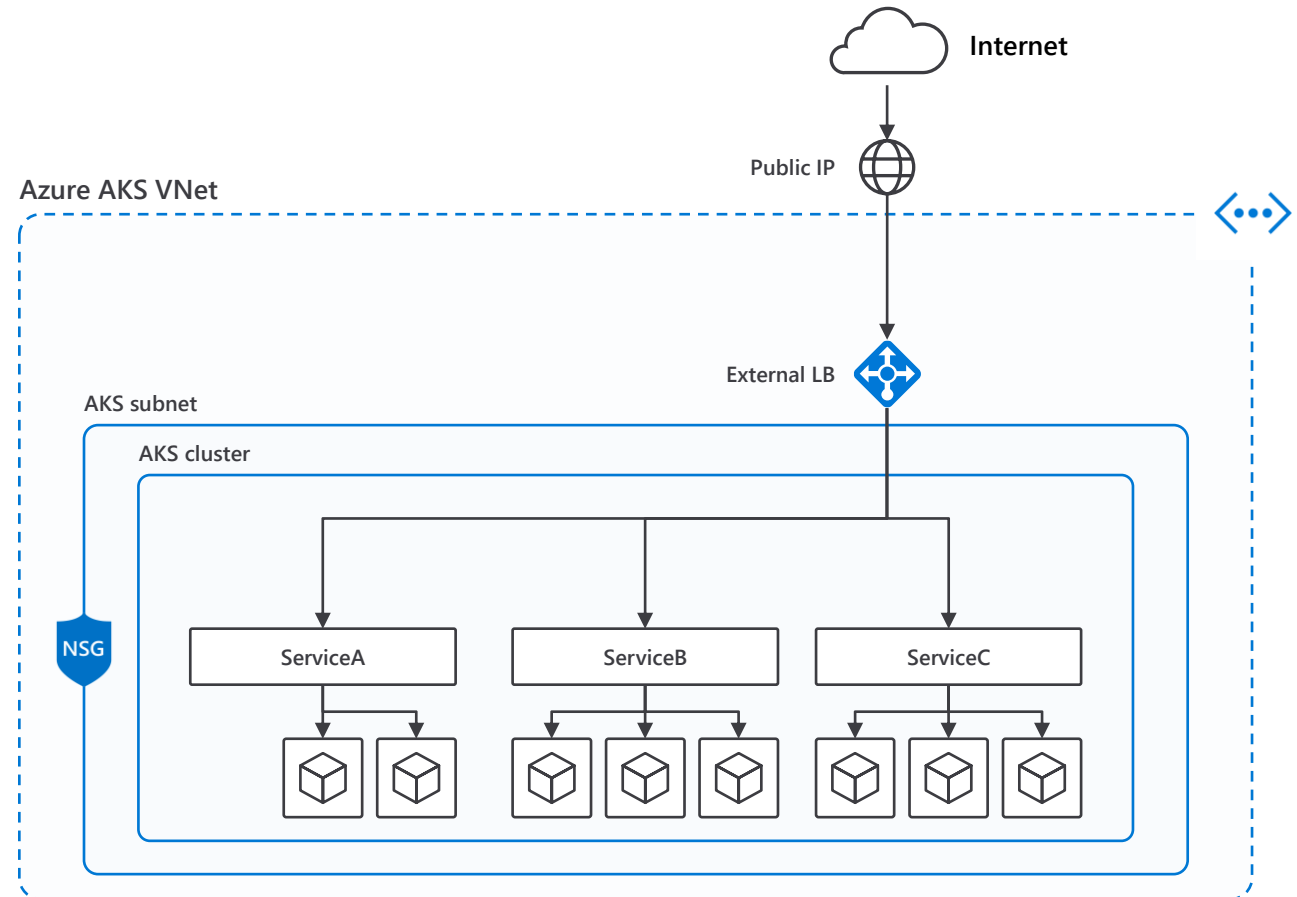
- Creates a port mapping on the underlying nodes that allows the application to be accessed directly with the node IP address and port



```
kind: Service
metadata:
  name: azure-vote-back
spec:
  type: NodePort
  ports:
    - port: 31000
  selector:
    app: azure-vote-back
```


# Services – Load balancer (external)

- Creates an Azure load balancer resource, configures a Public external IP address, and connects the requested pods to the load balancer backend pool
- To allow customers traffic to reach the application, load balancing rules are created on the desired ports



# Services – Load balancer (external)

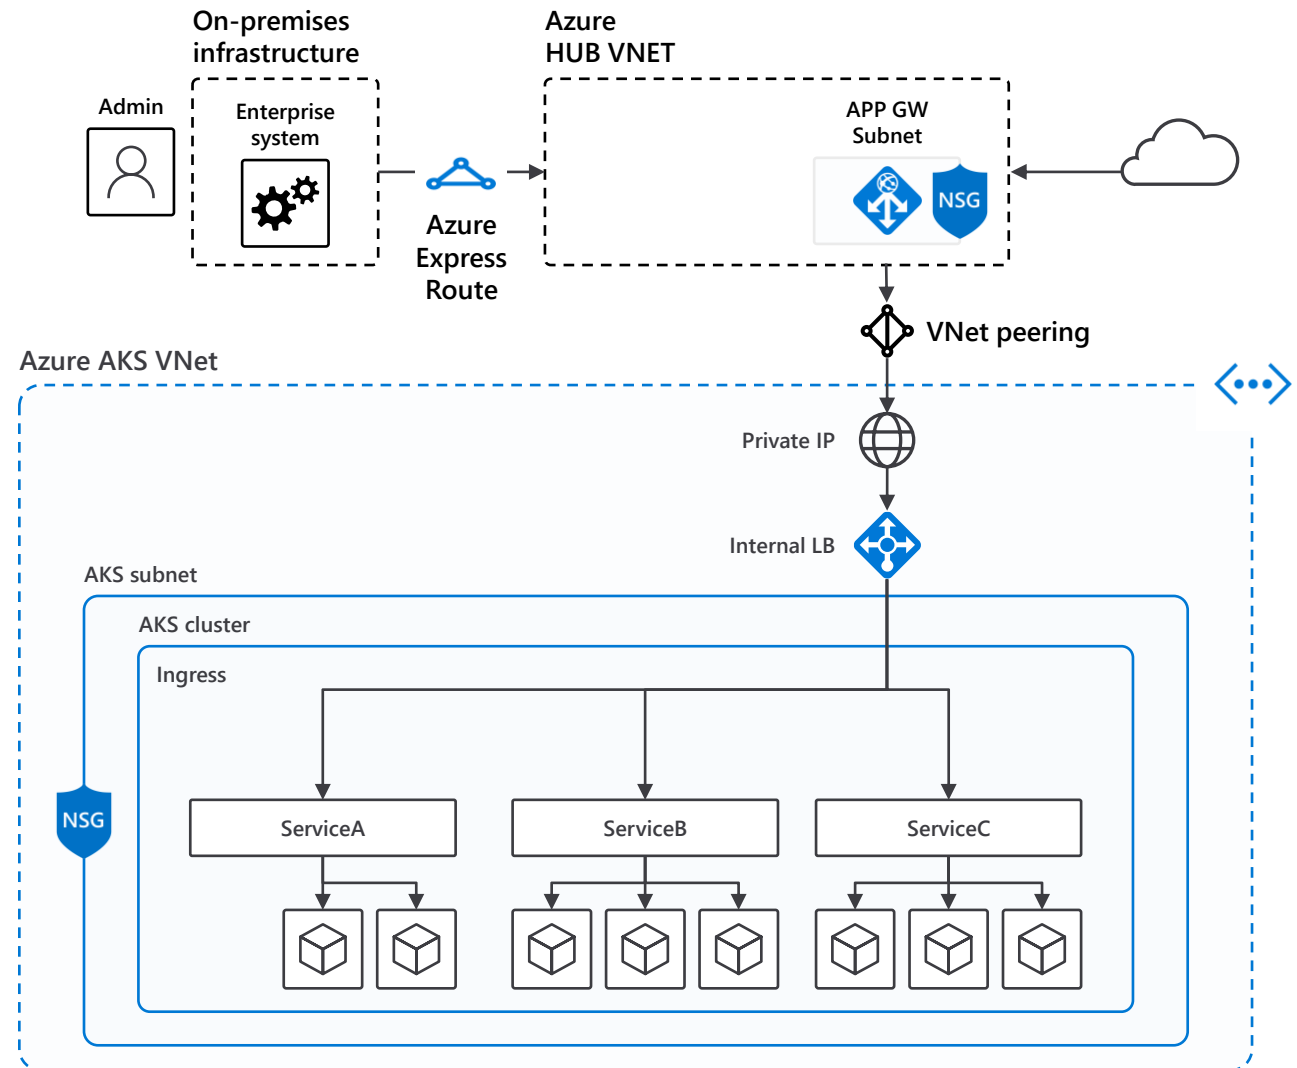
- Creates an Azure load balancer resource, configures a Public external IP address, and connects the requested pods to the load balancer backend pool
- To allow customers traffic to reach the application, load balancing rules are created on the desired ports



```
kind: Service
metadata:
  name: azure-vote-front
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: azure-vote-front
```


# Services – Load balancer (internal)

- Creates an Azure load balancer resource, configures a VNET external IP address, and connects the requested pods to the load balancer backend pool
- To allow customers traffic to reach the application, load balancing rules are created on the desired ports



# Services – Load balancer (internal)

- Creates an Azure load balancer resource, configures a VNET external IP address, and connects the requested pods to the load balancer backend pool
- To allow customers traffic to reach the application, load balancing rules are created on the desired ports



```
apiVersion: v1
kind: Service
metadata:
  name: azure-vote-front-internal
  annotations:
    service.beta.kubernetes.io/azure-load-balancer-internal: "true"
spec:
  type: LoadBalancer
  ports:
    - port: 80
  selector:
    app: azure-vote-front
```



# Lab 4: Services

## Expose different types of services



Task



With Azure

Expose public service

```
kubectl create -f azure-vote-service-public.yaml
```

```
# get public IP
```

```
kubectl get service azure-vote-front -watch
```

Expose private (Cluster IP)  
service

```
kubectl create -f azure-vote-service-private.yaml
```

```
# get IP
```

```
kubectl get services
```

```
# run test pod to for curl
```

```
kubectl run -i --tty --rm debug --image=radial/busyboxplus --restart=Never -- sh
```

Expose NodePort service

```
kubectl create -f nodeport-service.yaml
```

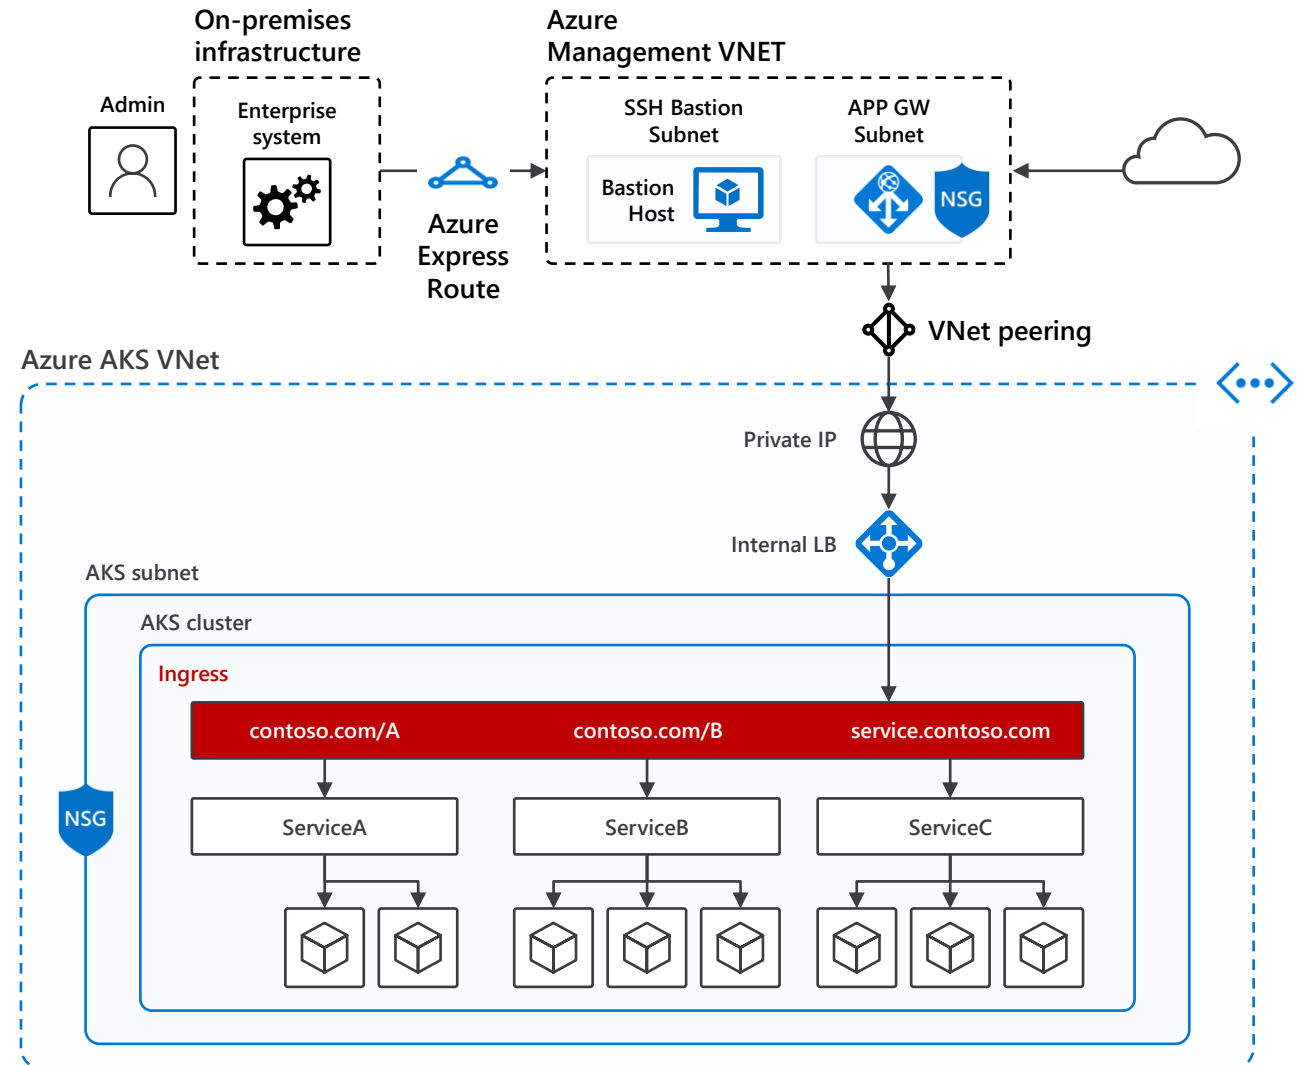
```
# get IP
```

```
kubectl describe service azure-vote-front-node
```

# Intelligent traffic distribution with Ingress Controllers

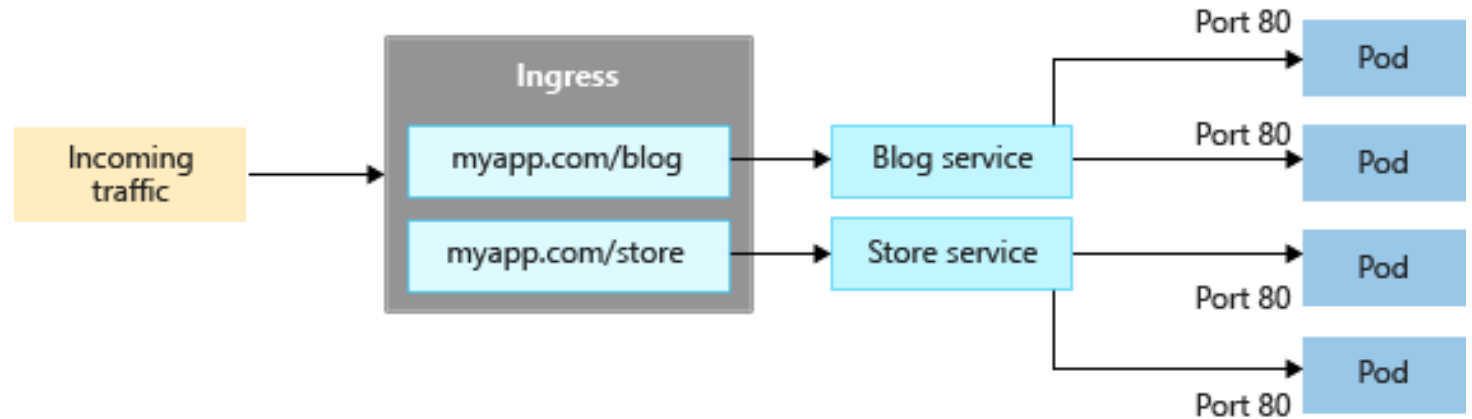
# Ingress controllers

- Ingress is an object that allows access to your Kubernetes services from outside the Kubernetes cluster. You configure access by creating a collection of rules that define which inbound connections reach which services.
- Ingress controllers work at layer 7 and can use more intelligent rules to distribute application traffic (e.g., route HTTP traffic to different applications based on the inbound URL).
- Typically, ingress controllers provide load balancing, SSL termination and name-based virtual hosting



# Ingress controllers

- A LoadBalancer type Service rely on an underlying Azure load balancer. The load balancer is configured to distribute traffic to the pods in your Service on a given port. The LoadBalancer only works at layer 4 - the Service is unaware of the actual applications and can't make any additional routing considerations.
- Ingress controllers work at layer 7 and can use more intelligent rules to distribute application traffic. A common use of an Ingress controller is to route HTTP traffic to different applications based on the inbound URL.



# Lab 5a: Setup Ingress Controller

Create basic ingress controller using nginx



## Task



## With Azure

Setup nginx as an ingress controller

```
# Create a namespace for your ingress resources
kubectl create namespace ingress-basic

# Add the official stable repo
helm repo add stable https://kubernetes-charts.storage.googleapis.com/

# Use Helm to deploy an NGINX ingress controller
helm install nginx stable/nginx-ingress \
  --namespace ingress-basic \
  --set controller.replicaCount=2 \
  --set controller.nodeSelector."beta\.kubernetes\.io/os"=linux \
  --set defaultBackend.nodeSelector."beta\.kubernetes\.io/os"=linux
```

Get ingress controller IP

```
# check-ingress-controller
kubectl get service -l app=nginx-ingress --namespace ingress-basic
```

# Lab 5b: Use Ingress Controllers

Expose APIs using ingress controllers



Task



With Azure

Create sample apps

```
kubectl create -f ingress/apple.yaml  
kubectl create -f ingress/banana.yaml
```

Define ingress rules

```
kubectl create -f ingress/app-ingress.yaml
```

Get ingress controller IP

```
kubectl get service -l app=nginx-ingress --namespace ingress-basic
```

Test it in your browser


Open your browser with the above IP and test:

```
http://<your-ingress-external-ip>/apple
```

```
http://<your-ingress-external-ip>/banana
```

# Containers & Kubernetes workshops

## Containers 101



## Kubernetes 101



## Kubernetes 201



Thank you