




Kubernetes on Azure 201 Workshop

Rui Félix Pereira

2020.05.05

Containers & Kubernetes workshops


Containers 101



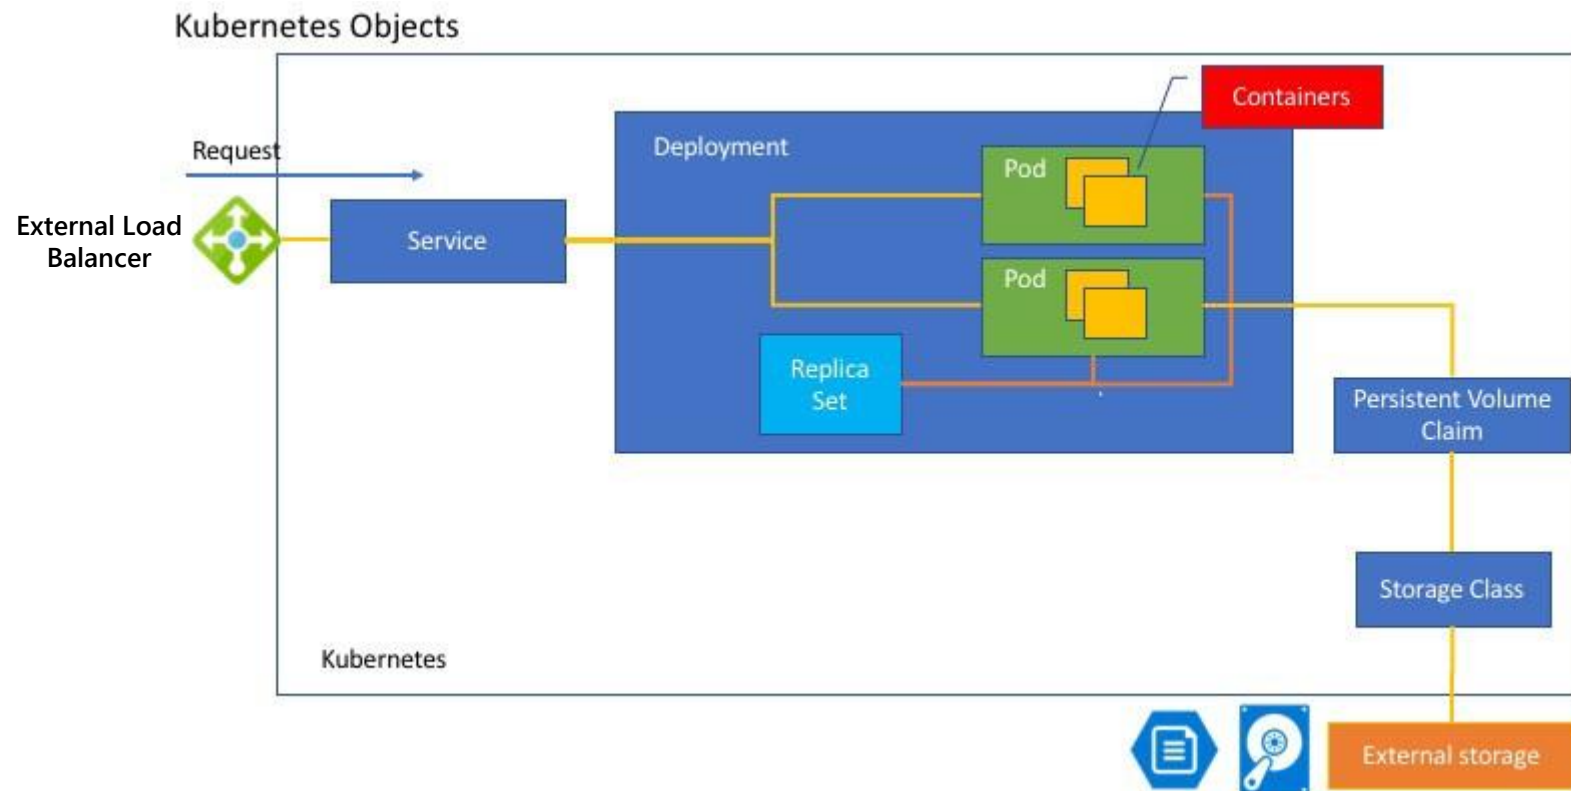
Kubernetes 101



Kubernetes 201



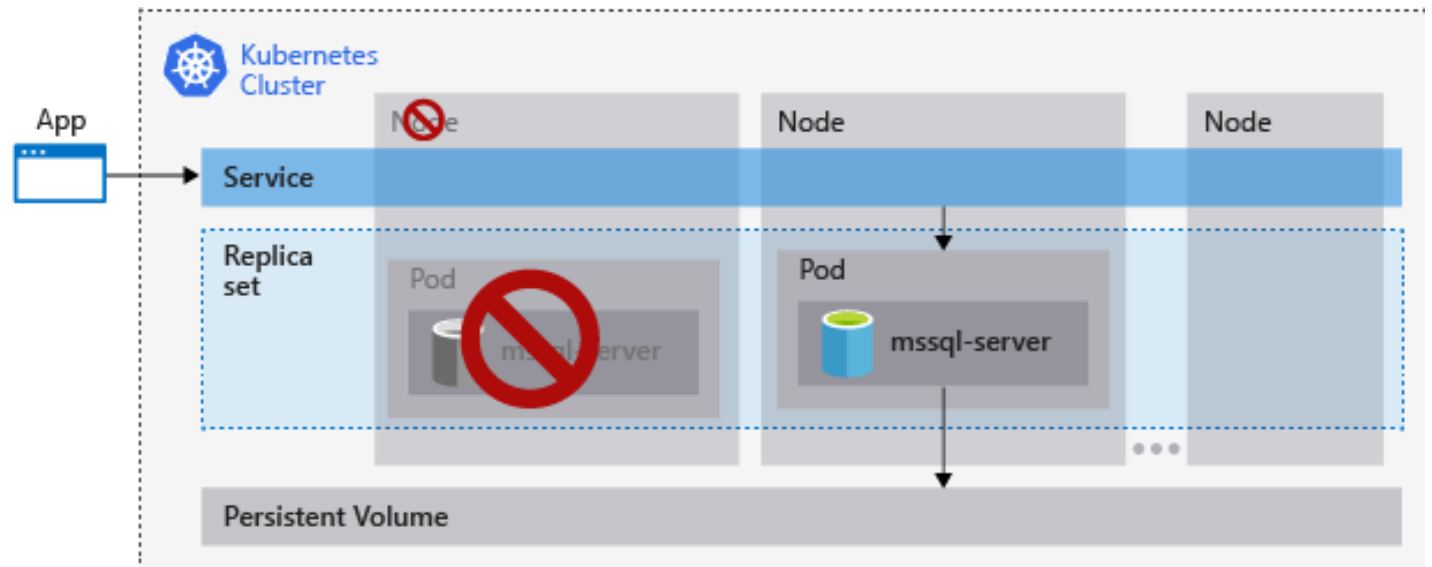
Kubernetes 101 recap



Using Persistent volumes

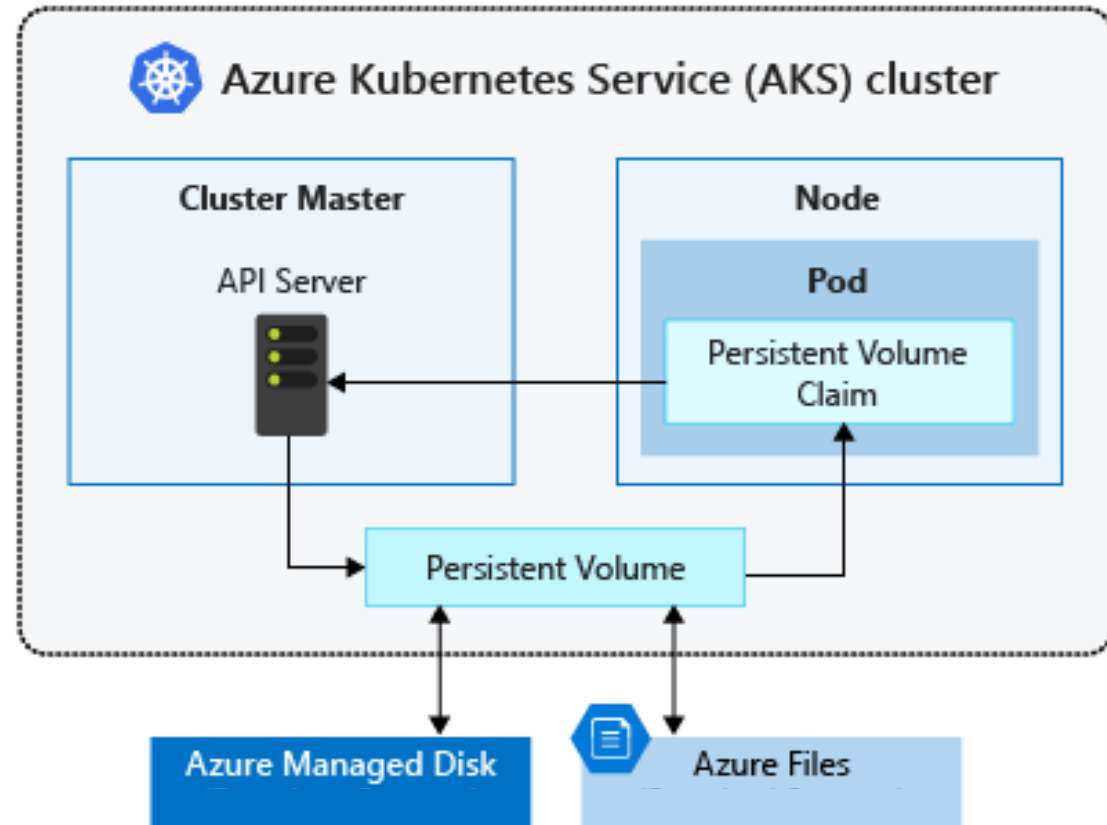
Storage and persistence

- Containers are volatile.
- In some cases, persistent storage is a need for sharing data between containers or to guarantee high availability where data must resist to container failures



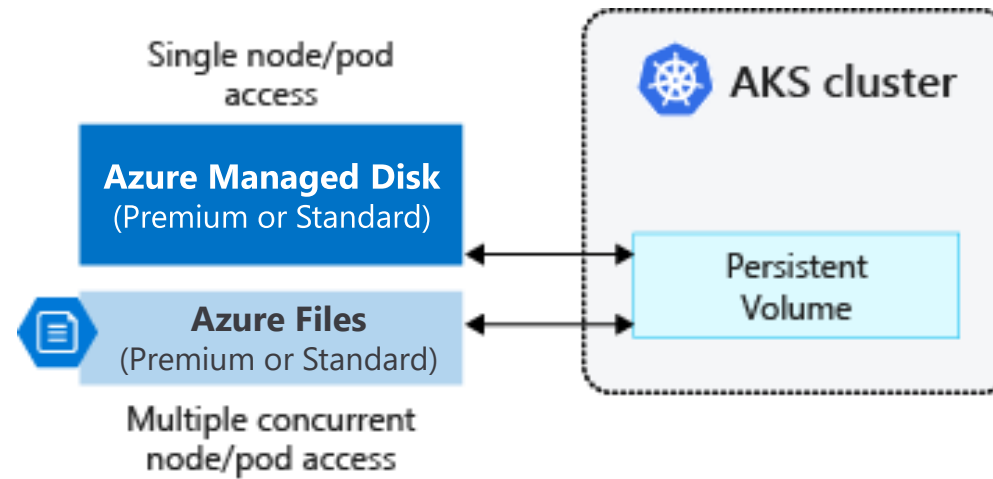
Storage and persistence on AKS

- A PersistentVolume can be **statically** created by a cluster administrator.
- Or **dynamically** created by the Kubernetes API server. If a pod is scheduled and requests storage that is not currently available, Kubernetes can create the underlying Azure Disk or Files storage and attach it to the pod. Dynamic provisioning uses a StorageClass to identify what type of Azure storage needs to be created.



Storage classes

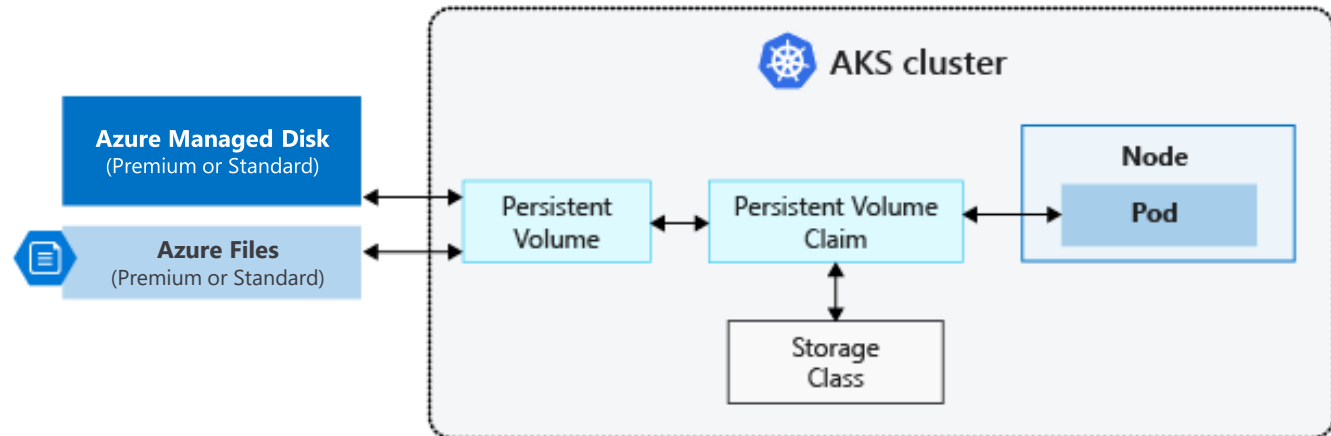
- A storage class is used to define how a unit of storage is dynamically created with a persistent volume
- Each AKS cluster includes 4 pre-created storage classes, both configured to work with Azure disks and files: Standard and Premium



```
kind: StorageClass
apiVersion: storage.k8s.io/v1beta1
metadata:
  name: azure-disk-standard
provisioner: kubernetes.io/azure-disk
parameters:
  storageaccounttype: Standard_LRS
  kind: Managed
```

Storage – Persistent volumes

- A persistent volume claim (PVC) is used to automatically provision storage based on a storage class



```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: my-data-pv-claim
  annotations:
    volume.beta.kubernetes.io/storage-class: default
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 2Gi
```


Storage – Using persistent volumes

- Using a persistent volume claim (PVC) in a Pod and mounting it in a mount point

```
kind: Pod
apiVersion: v1
metadata:
  name: task-pv-pod
spec:
  volumes:
    - name: task-pv-storage
      persistentVolumeClaim:
        claimName: my-data-pv-claim
  containers:
    - name: task-pv-container
      image: nginx
      ports:
        - containerPort: 80
          name: "http-server"
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: task-pv-storage
```

Lab 1: Dynamic Persistent Volumes

Create and use new persistent volume



Task



With Azure

Check storage classes

```
kubectl get storageclasses
```

```
# create new storage class if needed
```

```
kubectl create -f sample-storageclass.yaml
```

Create persistent volume

```
kubectl create -f sample-pvc.yaml
```

```
# get persistent volumes
```

```
kubectl get pv
```

```
kubectl describe pv
```

Use persistent volume

```
kubectl apply -f test-persistent-volumes.yaml
```

```
# get pods using persistent volume
```

```
kubectl get pod task-pv-pod
```

```
# access persistent volume (write something, destroy the pod and repeat the test again)
```

```
kubectl exec -it task-pv-pod -- /bin/bash
```

Static storage volumes not managed by AKS

- A PersistentVolume can be **statically** created by a cluster administrator.
- Create the Azure Disk or Azure File Share and collect the resource URI
- Configure it in a Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod-disk-tst
spec:
  containers:
  - image: nginx:1.15.5
    name: mypod
    volumeMounts:
    - name: azure
      mountPath: /mnt/azure
  volumes:
  - name: azure
    azureDisk:
      kind: Managed
      diskName: mydisk01
      diskURI: /subscriptions/2dd567af-d55f-4c2c-ba14-4bd7699a59b3/resourceGroups/k8s-volumes/providers/Microsoft.Compute/disks/mydisk01
```

Lab 2: Static Volumes

Use existing static volumes not managed by AKS



Task



With Azure

Create Azure Disk or Azure File Share and collect the Id

```
# Sample Azure Disk Id  
/subscriptions/2dd567af-d55f-4c2c-ba14-4bd7699a59b3/resourceGroups/k8s-volumes/providers/Microsoft.Compute/disks/mydisk01
```

Deploy Pod using the created disk

```
kubectl apply -f azure-disk-pod.yaml
```

Test and use the persistent volume inside the Pod


```
# access persistent volume (write something, destroy the pod and repeat the test again)  
kubectl exec -it mypod-disk-tst -- /bin/bash
```

Dynamic Provisioning

- Use Dynamic Provisioning whenever possible


StorageClass

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
...
parameters:
  kind: Managed
  storageaccounttype: Premium_LRS
provisioner: kubernetes.io/azure-disk
reclaimPolicy: Delete
volumeBindingMode: Immediate.
```



PersistentVolumeClaim

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: azure-managed-disk
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: managed-premium
resources:
  requests:
    storage: 5Gi
```



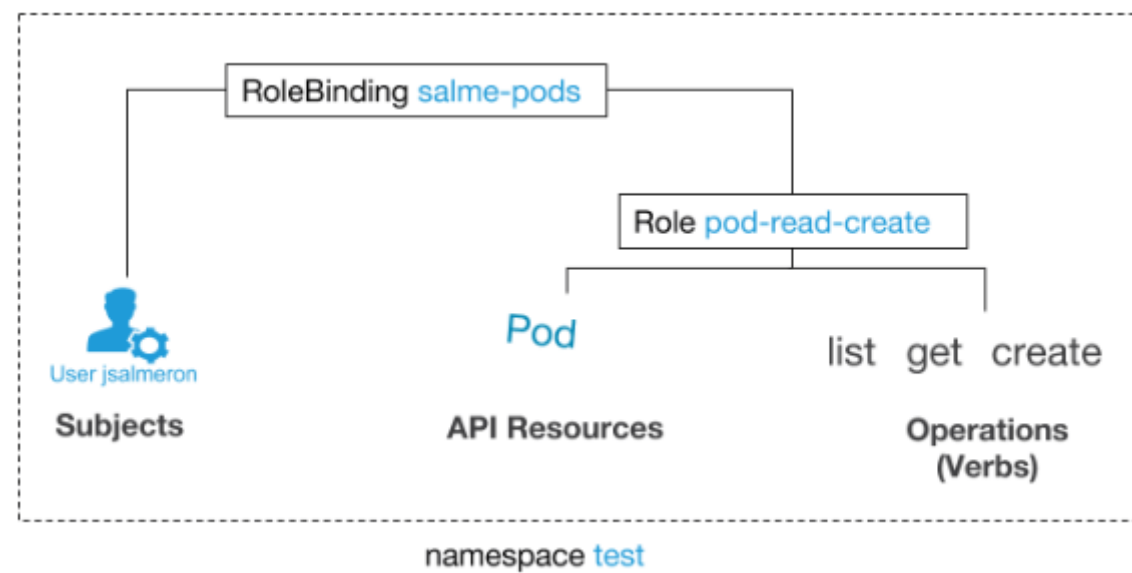
Pod

```
kind: Pod
..
spec:
  containers:
    - name: myfrontend
      image: nginx
      volumeMounts:
        - mountPath: "/mnt/azure"
          name: volume
  volumes:
    - name: volume
      persistentVolumeClaim:
        claimName: azure-managed-disk
```


Role based access control (RBAC)


Role based access control (RBAC)

- Role grants permissions to Kubernetes objects, typically in a namespaces
- RoleBinding can be assigned to users or groups



AKS RBAC with Azure Active Directory

- AKS uses Azure Active Directory as an Identity Services for users and groups



RBAC Role

- Roles grants permissions to Kubernetes objects, typically namespaces

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: dev-user-full-access
  namespace: dev
rules:
- apiGroups: ["", "extensions", "apps"]
  resources: ["*"]
  verbs: ["*"]
- apiGroups: ["batch"]
  resources:
  - jobs
  - cronjobs
  verbs: ["*"]
```

RBAC RoleBinding

- RoleBinding can be assigned to users or groups and optionally to namespaces

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: dev-user-access
  namespace: dev
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: Role
  name: dev-user-full-access
subjects:
- kind: Group
  namespace: dev
  name: groupObjectId
```


RBAC ClusterRoleBinding

- These are cluster roles that can be assigned to users and groups

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: contoso-cluster-admins
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: User
  name: "user@contoso.com"
```

Lab 3a: RBAC

Protect namespaces with RBAC



Task



With Azure

Create AKS with RBAC support

```
# This assumes that you have an AKS created with RBAC support  
https://docs.microsoft.com/en-us/azure/aks/azure-ad-integration-cli
```



```
# Get resource ID of the AKS in the Azure Active Directory  
AKS_ID=$(az aks show --resource-group myResourceGroup --name myAKSCluster --query id -o tsv)
```

Create users and groups in
Azure Active Directory

```
# Create group for Application developers: appdev group.  
APPDEV_ID=$(az ad group create --display-name appdev --mail-nickname appdev --query objectId -o tsv)  
  
# Assign it as a user of the AKS  
az role assignment create --assignee $APPDEV_ID --role "Azure Kubernetes Service Cluster User Role" \  
  --scope $AKS_ID  
  
# Create and Application developer user named aksdev that is part of the appdev group.  
AKSDEV_ID=$(az ad user create --display-name "AKS Dev" --user-principal-name aksdev@contoso.com \  
  --password P@ssw0rd1 --query objectId -o tsv)  
  
# Assign it as a member of the appdev group  
az ad group member add --group appdev --member-id $AKSDEV_ID
```

Lab 3b: RBAC


Protect namespaces with RBAC

 Task	 With Azure
Get admin credentials for AKs	<code>az aks get-credentials --resource-group myResourceGroup --name myAKSCluster -admin</code>
Create a new namespace	<code>kubectl create namespace dev</code>
Create a role for namespace dev	<code>kubectl apply -f role-dev-namespace.yaml</code>
Get the resource ID for the appdev group in the Azure Active Directory	<code>az ad group show --group appdev --query objectId -o tsv</code>
Create a RoleBinding for the appdev group and the previously created Role	<p># On the last line of file rolebinding-dev-namespace.yaml, replace <code>groupObjectId</code> with the group object ID output from the previous command</p> <code>kubectl apply -f rolebinding-dev-namespace.yaml</code>
Test it with a user	<code>az aks get-credentials --resource-group myResourceGroup --name myAKSCluster --overwrite-existing</code>

Manage configurations

ConfigMaps

- ConfigMaps are useful for storing and sharing non-sensitive, unencrypted configuration information.
- ConfigMaps bind configuration files, command-line arguments, environment variables, port numbers, and other configuration artifacts to your Pods' containers and system components at runtime
- Allow you to separate your configurations from your Pods and components, preventing hardcoding configuration data to Pod specifications



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-special-config
  namespace: default
data:
  log_level: INFO
  special.type: xpto
```


ConfigMaps data usage

- ConfigMaps data can be consumed in pods in a variety of ways
- Populate the values of environment variables
- Set command-line arguments in a container
- Populate config files in a volume

```
apiVersion: v1
kind: Pod
metadata:
  name: dapi-test-pod
spec:
  containers:
    - name: test-container
      image: busybox
      command: [ "/bin/sh", "-c", "env" ]
      env:
        - name: SPECIAL_TYPE_KEY
          valueFrom:
            configMapKeyRef:
              name: my-special-config
              key: special.type
  restartPolicy: Never
```

Lab 4: ConfigMap

Create and use ConfigMap from Pod



Task



With Azure

Create ConfigMap

```
# Create config map from poreperties files
kubectl create configmap myconfigmap --from-file xpto.properties

# Or create it form yaml
kubectl create -f myconfigmap.yaml
```

Get and test ConfigMap

```
# Get config map value
kubectl get configmap my-special-config -o yaml
```

Use the ConfigMap from Pod

```
kubectl create -f pod-using-configmap.yaml
```

Manage
secrets

Secrets

- Secrets are useful for storing and sharing sensitive, encrypted configuration information.
- Allow you to separate your sensitive secrets (e.g., passwords, connection strings, etc.) from your Pods, preventing hardcoding

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  username: YWRtaW4=
  password: MWYyZDF1MmU2N2Rm
```



Secrets usage

- Secrets, like ConfigMaps data can be consumed in pods in a variety of ways
- Populate the values of environment variables
- Set command-line arguments in a container
- Populate secret files in a volume

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
    - name: mycontainer
      image: redis
      env:
        - name: SECRET_USERNAME
          valueFrom:
            secretKeyRef:
              name: mysecret
              key: username
        - name: SECRET_PASSWORD
          valueFrom:
            secretKeyRef:
              name: mysecret
              key: password
      restartPolicy: Never
```


Lab 5: Secret

Create and use Secret from Pod

 Task	 With Azure
Create sample secrets	<pre>echo -n 'admin' base64 YWRtaW4= echo -n '1f2d1e2e67df' base64 MWYyZDFlMmU2N2Rm</pre>
Create Secret	<pre>kubectl create -f secret.yaml</pre>
Get and test Secret	<pre># Get secret value kubectl get secret mysecret -o yaml # Decode secret value echo 'MWYyZDFlMmU2N2Rm' base64 -decode</pre>
Use the Secret from Pod	<pre>kubectl create -f pod-using-secret.yaml</pre>

Kubernetes Operators

Custom Resources - Definition

- Kubernetes is highly extensible
- We can define Custom Resources on top of the out-of-the-box objects/resource types/APIs that come with k8s
- Defines a new API in K8S
- Seamless integration with existing APIs
- We can use kubectl

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: tomcats.tomcat.apache.org
spec:
  group: tomcat.apache.org
  names:
    kind: Tomcat
    listKind: TomcatList
    plural: tomcats
    singular: tomcat
  scope: Namespaced
  subresources:
    status: {}
  validation:
    openAPIV3Schema:
      ...
```

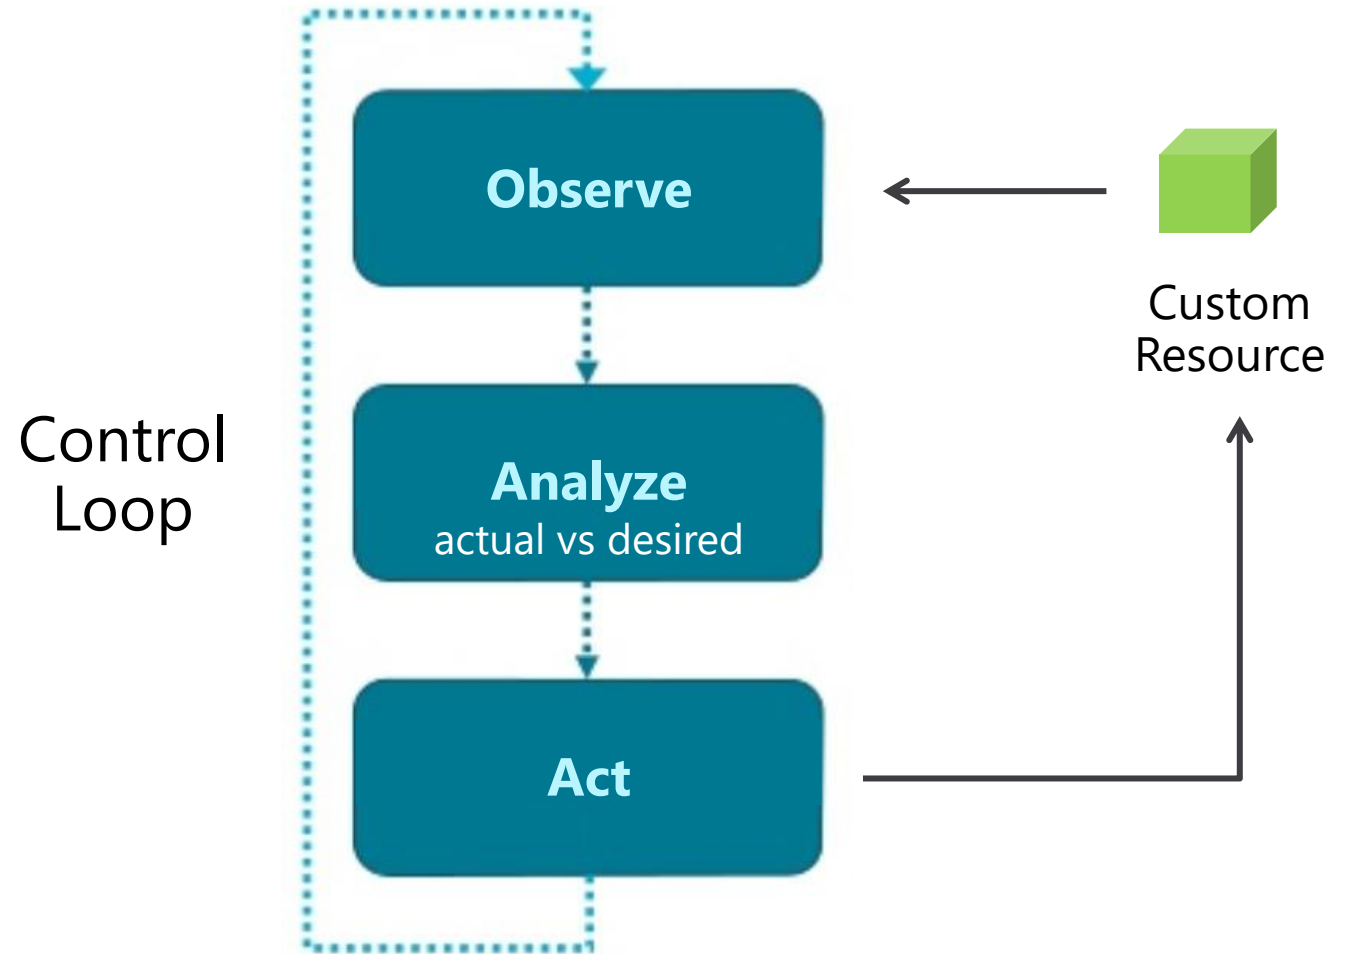
Custom Resources – Object/Instance

- We can define and configure K8S objects out of the custom resource definitions, just like any other Kubernetes object
- Using YAML and kubectl ...

```
apiVersion: tomcat.apache.org/v1alpha1
kind: Tomcat
metadata:
  name: tomcat
spec:
  replicas: 2
  image: tomcat:latest
  imagePullPolicy: IfNotPresent
  webArchiveImage: ananwaresystems/webarchive:1.0
  deployDirectory: /usr/local/tomcat/webapps
```

Custom Resources & Operators

- Remember that K8S uses the concept of desired state configuration
- We can automate the control of these custom resources using Operators
- Operator Watch CR objects
- Analyze differences between Actual and Desired State
- Act on changes
- An operator is itself a Deployment/Pod
- You can create your own operators



Operators, operators, operators

- An operator is itself a Deployment/Pod
- You can create your own operators
- A lot of operators exist today to automate the management of stateful or more complex resources like a database or big application
- Extend and automate the native K8S automation capabilities, targeting specific scenarios and workloads

Good examples

<https://github.com/coreos/etcd-operator>



<https://github.com/operator-framework/awesome-operators>



<https://operatorhub.io/>

Lab 6: Operators

Create a custom resource and deploy an operator to manage it

 Task	 With Azure
Create a custom resource	<pre>kubectl apply -f lab06-operator/tomcat-crd-definition.yaml</pre>
Create an operator for the custom resource	<pre># operator runs with a service account and a specific role kubectl apply -f lab06-operator/service_account.yaml kubectl apply -f lab06-operator/role.yaml kubectl apply -f lab06-operator/role_binding.yaml # deploy the operator kubectl apply -f lab06-operator/operator.yaml</pre>
Create an instance of the custom resource type	<pre># Create a Tomcat cluster kubectl apply -f lab06-operator/tomcat-cr-instance.yaml</pre>
Check what is deployed and running	<pre>kubectl get tomcat kubectl get pod kubectl get svc</pre>

Package apps with Helm

Helm

- Package manager for Kubernetes
- Helm helps you manage Kubernetes applications
- Avoiding K8S yml templates copy & paste all the time
- Helm Charts help you define, install, and upgrade even the most complex Kubernetes application
- Charts are easy to create, version, share, and publish — so start using Helm and stop the copy-and-paste.



<https://github.com/helm/charts>

Lab 7: Helm

Create and use a new Helm chart



Task



With Azure

Create a new Helm chart/package

```
helm create webfrontend
```

Edit your chart values

```
# Make the following updates to `webfrontend/values.yaml`:  
- Change the image of your container from `image.repository` to `microsoft/azure-vote-front:v1`  
- Change the service type to expose your application from `service.type` to `LoadBalancer`
```

```
## Edit your chart version  
Update `appVersion` to `v1` in `webfrontend/Chart.yaml`
```

Install your chart

```
helm install webfrontend webfrontend/
```


Validate it!!!

```
# Validate that your chart exists  
helm list
```


```
# Validate deployed objects  
kubectl get all
```

Containers & Kubernetes workshops


Containers 101



Kubernetes 101



Kubernetes 201



Thank you